

Commune: Shared Ownership in an Agnostic Cloud

Claudio Soriente
ETH Zurich, Switzerland
claudio.soriente@inf.ethz.ch

Ghassan O. Karame
NEC Laboratories Europe, Germany
ghassan.karame@neclab.eu

Hubert Ritzdorf
ETH Zurich, Switzerland
hubert.ritzdorf@inf.ethz.ch

Srdjan Marinovic
The Wireless Registry Inc, USA
srdjan@wirelessregistry.com

Srdjan Capkun
ETH Zurich, Switzerland
srdjan.capkun@inf.ethz.ch

ABSTRACT

Cloud storage platforms promise a convenient way for users to share files and engage in collaborations, yet they require all files to have a single owner who unilaterally makes access control decisions. Existing clouds are, thus, agnostic to shared ownership. This can be a significant limitation in many collaborations because, for example, one owner can delete files and revoke access without consulting the other collaborators.

In this paper, we first formally define a notion of *shared ownership* within a file access control model. We then propose a solution, called **Commune**, to the problem of distributed enforcement of shared ownership in agnostic clouds, so that access grants require the support of an agreed threshold of owners. **Commune** can be used in existing clouds without modifications to the platforms. We analyze the security of our solution and evaluate its performance through an implementation integrated with Amazon S3.

Categories and Subject Descriptors

C.2.0 [Computer-Communication Networks]: General – Security and protection.

Keywords

Cloud security; Shared ownership; Distributed enforcement

1. INTRODUCTION

Even though the cloud promises a convenient way for users to share files and effortlessly engage in collaborations, it still retains the notion of *individual* file ownership. That is, each file stored in the cloud is owned by a single user, who can *unilaterally* decide whether to grant or deny any access request to that file. However, the individual ownership is not suitable for numerous cloud-based applications and collaborations. Consider a scenario where a number of research organizations and industrial partners want to set

up a shared cloud repository to collaborate on a joint research project. If all participants contribute their research efforts to the project, then they may want to share the ownership over the collaboration files so that all access decisions are agreed upon among the owners. There are two main arguments why this may be preferred to individual ownership. First, a sole owner can abuse his rights by unilaterally making access control decisions. The community features a number of anecdotes where malicious users revoke access to shared files from other collaborators. Second, even if owners are willing to elect and trust one of them to make access control decisions, the elected owner may not want to be held accountable for collecting and correctly evaluating other owners' policies. For example, incorrect evaluations may incur negative reputation or financial penalties.

In contrast to individual ownership, we introduce a novel notion of *shared ownership* where n users jointly own a file and each file access request must be granted by a pre-arranged threshold of t owners. We remark that existing cloud platforms, such as Amazon S3 and Dropbox, provide no support for shared ownership policies, and offer only basic access control lists. In short, they are *agnostic* to the concept of shared ownership. Furthermore, state-of-the-art trust management systems that can support shared ownership policies (e.g., SecPAL [9], KeyNote [12], Delegation Logic [22]) make all access decisions using a *centralized* Policy Decision Point (PDP). This is not suitable for enforcing our shared ownership model, because the user who administers the PDP can arbitrarily change the policy rules set by the owners and enforce his own policies.

In this paper, we address the problem of *distributed enforcement of shared ownership within an agnostic cloud*. By distributed enforcement, we mean enforcement where access to files in a shared repository is granted if and only if t out of n owners separately support the grant decision. To tackle this problem, we first introduce the Shared-Ownership file access control Model (SOM) to define our notion of shared ownership, and to formally state the given enforcement problem. We then propose our solution, called **Commune**, that enforces shared ownership policies in a distributed fashion. **Commune** can be used within a third-party cloud without any modifications to the platform. It only requires that the cloud offers basic access control lists, as is the case with current platforms. We integrate a prototype implementation of **Commune** within Amazon S3 [1] and we show that its performance scales well with the file size and with the number of users.

Permission to make digital or hard copies of all or part of this work for personal or classroom use is granted without fee provided that copies are not made or distributed for profit or commercial advantage and that copies bear this notice and the full citation on the first page. Copyrights for components of this work owned by others than ACM must be honored. Abstracting with credit is permitted. To copy otherwise, or republish, to post on servers or to redistribute to lists, requires prior specific permission and/or a fee. Request permissions from permissions@acm.org.

SACMAT'15, June 1–3, 2015, Vienna, Austria.

Copyright © 2015 ACM 978-1-4503-3556-0/15/06 ...\$15.00.

<http://dx.doi.org/10.1145/2752952.2752972>.

To the best of our knowledge, **Commune** is the first solution to distributed enforcement of shared ownership in an agnostic cloud. We summarize our contributions as follows:

- We formalize the notion of shared ownership within a file access control model named **SOM**, and use it to define a novel access control problem of distributed enforcement of shared ownership in an agnostic cloud.
- We propose a solution, called **Commune**, which distributively enforces **SOM** and can be deployed in an agnostic cloud platform. **Commune** ensures that (i) a user cannot read a file from a shared repository unless that user is granted read access by at least t of the owners, and (ii) a user cannot write a file to a shared repository unless that user is granted write access by at least t of the owners.
- We build a prototype of **Commune** and evaluate it within Amazon S3. We show that our solution scales well with the file size and the number of users.

The rest of the paper is organized as follows. Section 2 introduces our notion of shared ownership in a file access control model. Section 3 details **Commune** and analyzes its security. Section 4 evaluates the performance of **Commune** through an implementation within Amazon S3. Section 5 reviews related work, and Section 6 provides concluding remarks.

2. SOM: SHARED-OWNERSHIP FILE ACCESS CONTROL MODEL

In this section, we first formalize our notion of Shared Ownership within a file access control Model named **SOM**. We then discuss the main shortcoming of centralized enforcement, and we define the problem of **SOM**'s distributed enforcement.

2.1 Syntax and Semantics

For simplicity, we do not consider directories (or other file groupings). A file is created with the following request

$$U \text{ reqs Create}(F, t, \mathcal{O}).$$

Upon receiving this request, **SOM** creates a file F , assigns a user U and all the users in \mathcal{O} as F 's owners, and sets the file's threshold to t . **SOM** grants requests for file creation to authenticated users if the new file name is unique.

To access a file, a user submits a request with an action he wishes to perform on the file

$$U \text{ reqs Action}(F).$$

SOM does not instantiate concrete file actions, these are left to implementations, and we use $\text{Action}(F)$ to denote a generic file access action on a file F .

If an owner O wishes to grant an action to U over F , then he issues a credential of the form

$$O \text{ says } U \text{ can Action}(F).$$

Intuitively, a credential is a certificate by an owner to support a user action. The full credential and request grammar is:

$$\begin{aligned} \text{credential} &::= u(i) \text{ says } u(j) \text{ can Action}(f) \\ \text{request} &::= u \text{ reqs Create}(f, t, o_1, \dots, o_n) \\ \text{request} &::= u \text{ reqs Action}(f) \\ u, f, o &::= \text{String} \\ t &::= \mathbb{N} \end{aligned}$$

File access requests are granted if and only if t out of n owners issue the corresponding credentials. For example, if the threshold for F is 2, then U can perform $\text{Action}(F)$ if the following credentials are present: $O \text{ says } U \text{ can Action}(F)$ and $O' \text{ says } U \text{ can Action}(F)$, where O and O' are two of F 's owners.

The **SOM** access control policy, Π_{SOM} , is a mapping from the set of all requests \mathcal{R} and the set of all credentials \mathcal{C} into the $\{\text{grant}, \text{deny}\}$ decision set, such that the file threshold is always respected.

Definition 1 (Shared Ownership Access Control Policy). *The **SOM** access control policy, denoted Π_{SOM} , is a mapping $\mathcal{R} \times \mathcal{C} \rightarrow \{\text{grant}, \text{deny}\}$ such that:*

$$\begin{aligned} \Pi_{\text{SOM}}(U \text{ reqs Action}(F), \text{Creds}) &\mapsto \text{grant} \text{ iff} \\ \{ (O_1 \text{ says } U \text{ can Action}(F)), \dots, \\ &(O_t \text{ says } U \text{ can Action}(F)) \} \subseteq \text{Creds}, \\ O_1 \neq \dots \neq O_t \text{ and } O_1 \in \mathcal{O}, \dots, O_t \in \mathcal{O}, \end{aligned}$$

where t is F 's threshold, and \mathcal{O} is the set of F 's owners.

We say that **SOM** grants a file access to a given request r and its accompanying credentials C if and only if $\Pi_{\text{SOM}}(r, C) = \text{grant}$. Note that by Definition 1, **SOM** treats owners only as sources of credentials, and does not implicitly grant them any additional access rights.

2.2 Centralized vs. Distributed Enforcement

Existing credential-based access control systems have the following enforcement model. A Policy Enforcement Point (PEP) has one designated Policy Decision Point (PDP), which collects all the required credentials and evaluates an access control policy (such as Π_{SOM}) for a given request. The PDP has one user who can administrate its access control policy. We refer to this enforcement model as *centralized* since a single policy decision point grants all access requests. Note that even if a PEP used multiple PDP components (managed by different users), it would still require an additional PDP to centrally decide how these decisions are combined. In Appendix A, we show how **SOM** can be centrally enforced. We do this by specifying the **SOM** access control policy as a Datalog logic program. A Datalog interpreter can then act as the PDP. Using Datalog interpreters as PDPs is also common in state-of-the-art access control systems (e.g., [9], [22]).

The key shortcoming of centralized enforcement is that the PDP's access decision is solely sufficient for granting access requests. The administrator, responsible for managing the rules stored at the PDP, can therefore change the policy rules to grant requests that lack the necessary credentials. In practice, this means that the shared ownership policy agreed upon by the owners can always be bypassed, and thus the notion of shared ownership nullified. The natural question to consider is how to enforce the **SOM** policy so that the agreed thresholds cannot be bypassed. We refer to such an enforcement solution as *distributed* because it must grant access if and only if t owners separately support the grant decision.

When considering a cloud as a collaboration platform, this enforcement issue is even more difficult because a cloud platform does not allow deployment of additional enforcement components. A cloud platform only supports basic access control policies via Access Control Lists (ACLs). We frame these concerns as the **SOM** distributed enforcement problem.

Enforcement	Ownership	
	Individual	Shared
Centralized	ACLs	Datalog-based Systems
Distributed	—	Commune

Table 1: Solutions for ownership enforcement in a cloud. Note that Datalog-based systems require modifications to the cloud platform.

Problem Statement: How can the SOM access control policy be distributively enforced within a cloud platform that supports only ACL-based PDPs?

We summarize the state-of-the-art with respect to this problem in Table 1.

3. COMMUNE: DISTRIBUTED ENFORCEMENT OF SHARED OWNERSHIP

This section presents Commune, a solution for distributed enforcement of the SOM access control policy in an agnostic cloud. As SOM does not specify concrete file access operations, we instantiate Commune with write and read actions. Before introducing our solution, we outline our cloud and attacker model.

3.1 Cloud and Attacker Model

We focus on a cloud storage platform, \mathcal{S} , where a set of users \mathcal{U} have personal accounts onto which they upload files. For example, users might set up their own personal clouds [5, 6], or might create personal accounts in existing public clouds. A user $U \in \mathcal{U}$ can unilaterally decide who has access to files stored on his account. In particular, \mathcal{S} allows each user to define access control policies of the type $p : \mathcal{U} \times \{\text{write}, \text{read}\} \rightarrow \{\text{grant}, \text{deny}\}$. We also assume that \mathcal{S} correctly enforces individual access control policies. This model reflects the functionalities provided by existing cloud platforms, such as Amazon S3.

Since we assume that \mathcal{S} authenticates users, we only focus on internal adversaries. An adversary may try to gain read access to a file even if fewer than t owners have issued the corresponding credentials. We refer to this adversary as a “malicious reader”. Alternatively, an adversary, who has been granted write access by fewer than t owners, may try to publish a file F as if F were authored by a user had been granted write access by t or more owners. We refer to this adversary as a “malicious writer”. We also consider sets of users who collude to escalate their access rights.

3.2 Overview of Commune

Before describing Commune, we make the following observations:

Observation 1. Commune’s files cannot be stored on a single user account.

Following the discussion regarding the centralized enforcement, a single user must not be charged with making unilateral grant and deny decisions. Otherwise, that user may abuse his rights and take unilateral access control decisions. A naïve solution where a file is encrypted (e.g., using a key shared among the owners) and the ciphertext is stored on a single account, allows that account holder to, for example,

unilaterally deny read access to the ciphertext. If the ciphertext cannot be read, any mechanism to distribute or recover the encryption key is of no help. We argue, therefore, that Commune cannot use a centralized repository because the repository owner can unilaterally grant or deny access to the files stored therein. Our alternative is to use a “shared repository”, which is an abstraction built on top of the owners’ personal accounts on \mathcal{S} .

Observation 2. Commune cannot support in-place writing.

If Commune were to allow in-place writing, then users who are granted write access could overwrite a file with “garbage”. This would equate to granting users the right to unilaterally delete the file, thus nullifying our efforts to prevent such scenarios. A standard alternative to in-place writing is to introduce “copy-on-write” mechanisms whereby a new file is created upon each file write operation. To optimize performance, Commune implements versioning and splits files into *units* (i.e., the unit of granularity of versioning) so that writing a new version of an existing file, only requires updating the units that have changed with respect to the previous version.

Observation 3. Commune cannot prevent users from disseminating a file through an out-of-band channel.

Access control solutions cannot prevent a user from distributing content through an out-of-band channel. For example, a user who rightfully reads a file can leak it to third parties. Similarly, a malicious writer can write a file and disseminate it through an out-of-band channel. For example, a user can publish files on his account on \mathcal{S} and make them available for others to read. We cannot prevent such behaviour. Commune, however, must at least allow honest readers, who abide to the protocol specification, to distinguish between the content written by malicious writers and the content written by honest writers.

Given these observations, Commune unfolds as follows. At system setup, users define the set of n owners \mathcal{O} and the threshold t (with $t \leq n$).¹ Commune abstracts the storage space of the owners’ accounts on \mathcal{S} as the “shared repository”. Each owner grants/denies read and write access on his account to users (including other owners) according to his individual access control policy. The distributed enforcement of the SOM access control policy then follows from the enforcement of the individual access policies set by each owner.

To write a file to the shared repository, the writer encodes the file in *tokens* and distributes the tokens to the owners’ accounts. A file is written to the shared repository if and only if the writer successfully distributes the file’s tokens onto at least t owners’ accounts. That is, a user has write access granted to the shared repository if and only if he has write access granted to at least t of the owners’ accounts. We refer to such a user as an “authorized writer”.

To read a file from the shared repository, the reader must fetch the file’s tokens from at least t distinct owners’ accounts. Therefore, a user has read access granted to the file if and only if he has read access granted to the file’s

¹Mechanisms to select the set of owners and the threshold t are outside of our scope. In settings like scientific collaboration scenarios, owners and thresholds are agreed by the partners.

Algorithm 1 AON-FFT(K, f_1, \dots, f_m)

```
1: Parse  $f_1, \dots, f_m$  as  $f_1^0, \dots, f_m^0$ 
2: for  $r \leftarrow 1$  to  $\log_2 m$  do ▷ round counter
3:   for  $i \leftarrow 0$  to  $\frac{m}{2^r} - 1$  do
4:     for  $j \leftarrow 1$  to  $2^{r-1}$  do
5:        $f_{j+i \cdot 2^r}^r || f_{j+i \cdot 2^r + 2^{r-1}}^r \leftarrow E(K, f_{j+i \cdot 2^r}^{r-1}, f_{j+i \cdot 2^r + 2^{r-1}}^{r-1})$ 
6:     end for
7:   end for
8: end for
9: return  $f_1^r, \dots, f_m^r$  as  $\bar{f}_1, \dots, \bar{f}_m$ 
```

tokens by at least t owners. We refer to such a user as an “authorized reader”.

To securely enforce shared ownership policies, **Commune** is designed to fulfil the following properties.

- **P1:** A malicious writer (i.e., a user who has been granted write access by fewer than t owners), must not be able to publish a file F as if F were authored by an authorized writer.
 - **P2:** A malicious reader (i.e., a user who has been granted read access to a file F by fewer than t owners), must not be able to recover the file content. This property must also hold in case of *revocation*. Assume that, at the time τ_1 , U has read access to F granted by at least t owners. Also assume that, at the time $\tau_2 > \tau_1$, U has his access rights revoked. This happens if, at the time τ_2 , some of the owners decide to revoke read access to U so that U is left with fewer than t read grants. We must ensure that, starting from time τ_2 , U cannot recover meaningful bits of F . We remark that, as is common for access control systems, we cannot prevent U from storing a local copy of F at the time t_1 and reading it even after his read right has been revoked.
- Commune** must also provide *collusion resistance*. That is, coalitions of users—where no single user is an authorized reader—must not be able to pool their credentials to escalate their read access rights.

Property P1 ensures protection against malicious writers who try to disseminate content despite lacking the required credentials. Property P2 guarantees that malicious readers cannot read content written to the shared repository.

Commune fulfils property P1 by design, through the abstraction of the shared repository and the copy-on-write mechanism (see Section 3.5). Property P2 is fulfilled through two cryptographic building blocks: Secure File Dispersal (SFD), and Collusion Resistant Secret Sharing (CRSS). SFD ensures that malicious readers cannot acquire any information about a file, even if they previously had access to the file and were later revoked. CRSS builds atop SFD and ensures that coalitions of users where no single user has enough credentials to read the file, cannot pool their credentials in order to escalate their read access rights.

In the following, we describe and analyze SFD (Section 3.3) and CRSS (Section 3.4). In Section 3.5, we detail the integration of both building blocks in **Commune**.

3.3 Secure File Dispersal (SFD)

Information dispersal algorithms [24] encode a file in n chunks so that any t chunks (where $t \leq n$) are sufficient to decode it. However, information dispersal algorithms do not provide any security guarantees if the number of available

chunks is smaller than t : any party with fewer than t chunks may still recover meaningful information about the original file’s content.

Previous work on securing information dispersal algorithms [25] combines erasure codes with All-Or-Nothing Transforms (AONT) [26]. The latter is an efficient block-wise transformation that maps an n -block bitstring in input to an n' -block bitstring in output (with $n' \geq n$). AONTs are designed in such a way that, unless all the n' output blocks are available, it is hard to recover any of the input blocks.

Existing AONTs [13, 26] leverage block ciphers and rely on the secrecy of a cryptographic key that is embedded within the output blocks. Given all AONT output blocks, the key can be recovered; once the key is known, individual blocks can be reverted, independently of other blocks. Current AONTs, therefore, preserve their all-or-nothing property only for *one time*: knowledge of the cryptographic key allows to revert single output blocks and to recover parts of the original data. This is at odds with our security requirements. As argued before, we cannot prevent users from caching a local copy of the file and reading it at later time when their read rights may have been revoked. However, we still want to provide revocation of a user who only stored the encryption key at the time when he had read access to the file.

We therefore introduce a new scheme, called Secure File Dispersal (SFD), that combines information dispersal algorithms with an AONT that preserves its all-or-nothing property even if the adversary has the encryption key.

Definition. An SFD scheme consists of the following algorithms:

$\{c_1, \dots, c_n\} \leftarrow \text{SFD.Encode}(t, n, F, K, \lambda)$. Encodes a file F into n chunks, such that F can be correctly decoded using any t chunks; K denotes a key used in the encoding process and λ is a security parameter.

$F' \leftarrow \text{SFD.Decode}(K, \mathcal{C}, \lambda)$. Takes as input a key K , a set of chunks \mathcal{C} , and security parameter λ ; it outputs a file F' .

Correctness. Given $\{c_1, \dots, c_n\} \leftarrow \text{SFD.Encode}(t, n, F, K, \lambda)$ and $F' \leftarrow \text{SFD.Decode}(K, \mathcal{C}, \lambda)$, we require that if $\mathcal{C} \subseteq \{c_1, \dots, c_n\}$ and $|\mathcal{C}| \geq t$, then $F' = F$.

Security. We define the advantage of adversary \mathcal{A} as follows:

$$\begin{aligned} \text{Adv}_{\text{SFD}}(\mathcal{A}) &= \Pr[f \leftarrow \mathcal{A}(K, \mathcal{C}) | K \leftarrow \{0, 1\}^l, l \geq \lambda, \\ &F = f_1, \dots, f_m \leftarrow \{0, 1\}^{m\lambda}, \\ &\{c_1, \dots, c_n\} \leftarrow \text{SFD.Encode}(t, n, F, K, \lambda), \\ &\mathcal{C} \subset \{c_1, \dots, c_n\}, |\mathcal{C}| < t, f \subseteq F, |f| \geq \lambda]. \end{aligned}$$

where $f \subseteq F$ refers to a substring of F . We say that SFD is secure if, for any p.p.t. adversary, its advantage is negligible in the security parameter, i.e., $\text{Adv}_{\text{SFD}}(\mathcal{A}) \leq \text{negl}(\lambda)$. Our security definition captures the scenario where, at an earlier time, \mathcal{A} was given enough chunks to decode F and has cached a copy of the key K , while at current time he is only given fewer than t chunks. Even if \mathcal{A} has the key K , we require the probability that \mathcal{A} recovers any λ consecutive bits of F to be negligible in the security parameter.




Figure 1: Sketch of the AON-FFT scheme where the input consists of $m = 8$ input blocks. Solid circles refer to the block cipher $E(\cdot)$, while empty circles depict its input/output blocks.

Instantiation. Our SFD scheme combines information dispersal techniques with AON-FFT, an all-or-nothing transformation inspired by Fast Fourier Transform.

Let $E : \{0, 1\}^{4\lambda} \rightarrow \{0, 1\}^{2\lambda}$ be a semantically secure block cipher (e.g., $E(\cdot)$ could correspond to 256-bit Rijndael [17], with $\lambda = 128$).² AON-FFT takes as input a symmetric key K (of size 2λ) and m input blocks f_1, \dots, f_m (each of size λ). It executes in $\log_2 m$ rounds and, at each round, applies $E(\cdot)$ to pairs of blocks. Each round is fed with the output of the previous round. The original input f_1, \dots, f_m is treated as the output of round 0; the final output of the algorithm is the output of round $\log_2 m$ (cf. Figure 1). The pseudo-code of AON-FFT is shown in Algorithm 1. We omit the details of the decryption algorithm since it is specular to encryption.

Given the pseudo-code of AON-FFT, our SFD scheme unfolds as follows:

$c_1, \dots, c_n \leftarrow \text{SFD.Encode}(t, n, F, K, \lambda)$. Parse F as f_1, \dots, f_m where each f_i has size λ .
Run $\tilde{f}_1, \dots, \tilde{f}_m \leftarrow \text{AON-FFT}(K, f_1, \dots, f_m)$. Use the information dispersal encoder to encode $\tilde{f}_1, \dots, \tilde{f}_m$ in n chunks with reconstruction threshold t .³

$F' \leftarrow \text{SFD.Decode}(K, C, \lambda)$. Given a set of at least t chunks C and key K , use the information dispersal decoder to decode blocks $\tilde{f}'_1, \dots, \tilde{f}'_m$. Run $f'_1, \dots, f'_m \leftarrow \text{AON-FFT}(K, \tilde{f}'_1, \dots, \tilde{f}'_m)$.

Correctness. If $\{c_1, \dots, c_n\} \leftarrow \text{SFD.Encode}(t, n, F, K, \lambda)$, any subset of at least t chunks $\{c_{i_1}, \dots, c_{i_t}\}$ can be decoded into the whole output of AON-FFT, namely $\tilde{f}_1, \dots, \tilde{f}_m$. Given K , the output of AON-FFT can be decrypted to recover $F = f_1, \dots, f_m$.

Security. Given the construction of our AON-FFT scheme, it is easy to see that each input block depends on all output blocks and on the encryption key. Furthermore, assuming that $E(\cdot)$ is a semantically secure block cipher, for any p.p.t. algorithm \mathcal{A} , we have $\text{Adv}_{\text{SFD}}(\mathcal{A}) \leq \text{negl}(\lambda)$. We provide a security argument in the Appendix.

Note that a construct similar to AON-FFT, was first mentioned by Rivest [26] and later on used as a “proof of storage”

²The key size is 2λ and the input/output size is also 2λ , totalling 4λ size of input.

³SFD can leverage any information dispersal algorithm (e.g., Reed-Solomon codes [30]).

in [29]. Nevertheless, the construction proposed therein can use any pseudo-random permutation in the FFT network. Our AON-FFT requires a keyed permutation, hence a block-cipher. Furthermore, the goal of the adversary in [29] is to recover, in a given amount of time, all *output* blocks. In contrast, the goal of our adversary is to recover any *input* block. This entails different security definition and analysis.

3.4 Collusion Resistant Secret Sharing (CRSS)

We now introduce our second building block, called Collusion Resistant Secret Sharing (CRSS). Similar to threshold secret-sharing schemes, CRSS allows one party to distribute a secret among a set of designated shareholders, so that any subset of shareholders of size equal to or greater than the threshold can reconstruct the secret. Furthermore, CRSS allows shareholders to issue to other users *delegation* to reconstruct the secret. If a user collects enough (i.e., above the threshold) delegations, he can rightfully reconstruct the secret. However, users cannot pool their delegations to reconstruct the secret, unless one of them has collected enough delegations. In *Commune*, CRSS is used to secret-share the key K used in SFD, in order to achieve collusion resistance.

CRSS is inspired by decentralized Attribute Based Encryption [21] where shares of a secret are *blinded* with shares of 0, such that, if a user collects enough shares for his identity, the blinding cancels out and the secret can be reconstructed.

Definition. Our definition of CRSS builds on top of a *standard* threshold secret-sharing scheme SS with algorithms $\text{SS.Share}(\cdot)$ and $\text{SS.Combine}(\cdot)$, to share and reconstruct a secret, respectively. We assume SS to be secure according to the **Game Priv** definition by Rogaway et al. [27]. That is, we assume that an adversary has only negligible advantage in identifying which out of two values was (t, n) secret-shared using the $\text{SS.Share}(\cdot)$ algorithm, even if the adversary can corrupt up to $t - 1$ shareholders and access their shares.

CRSS defines the following algorithms:

$\{s_1, \dots, s_n\} \leftarrow \text{CRSS.Share}(s, t, n)$. Shares secret s in a set of n shares $\{s_1, \dots, s_n\}$ with reconstruction threshold t .

$d_{i,j} \leftarrow \text{CRSS.Delegate}(s_i, U_j)$. Takes as input a share s_i and an user identity U_j . The output is a *delegation* $d_{i,j}$.

$s' \leftarrow \text{CRSS.Combine}(\{d_{i_1,j}, \dots, d_{i_l,j}\})$. Combines delegations $\{d_{i_1,j}, \dots, d_{i_l,j}\}$ into s' .

Correctness. Given $\{s_1, \dots, s_n\} \leftarrow \text{CRSS.Share}(s, t, n)$ and $s' \leftarrow \text{CRSS.Combine}(\{d_{i_1,j}, \dots, d_{i_l,j}\})$, we require that if $d_{i_p,j} \leftarrow \text{CRSS.Delegate}(s_{i_p}, U_j)$, for $1 \leq p \leq l$ and $l \geq t$, then $s' = s$.

Security. We model the security of CRSS using an adaptation of the **Game Priv** of [27] and we denote the refined game by **Game Priv***:

Init. The adversary \mathcal{A} submits two messages x_0, x_1 of equal length. The challenger flips an unbiased coin b and runs $\{s_1, \dots, s_n\} \leftarrow \text{CRSS.Share}(x_b, t, n)$.

Find. \mathcal{A} can submit two types of queries. In Type-1 queries, the adversary can corrupt up to $t' \leq t - 1$ shareholders and receives their shares. At this time, \mathcal{A}

picks t' indexes $i_1, \dots, i_{t'}$ and receives $\{s_{i_1}, \dots, s_{i_{t'}}\}$. In Type-2 queries, for any fresh identity U_j , the adversary can ask for up to t'' delegations, as long as $t' + t'' \leq t - 1$. \mathcal{A} submits an identity U_j and t'' indexes $i_1, \dots, i_{t''}$, and receives delegations $\{d_{i_1,j}, \dots, d_{i_{t''},j}\}$.

Guess. The adversary outputs his guess b' and wins if $b' = b$.

We define the advantage of the adversary as the probability of its winning minus a half. That is, $\text{Adv}_{\text{CRSS}}^{\text{Priv}^*}(\mathcal{A}) = \text{Prob}[\text{Priv}^*^{\mathcal{A}}] - \frac{1}{2}$. Therefore, we say that CRSS is secure if any p.p.t. algorithm \mathcal{A} has only negligible advantage in winning **Game Priv***.

The above **Game Priv*** models a scenario where a set of malicious users, including up to t' shareholders, collects up to t'' delegations for each of their identities. If $t' + t'' \geq t$, the malicious shareholders can produce the missing delegations for any of the colluding user identities, so that the secret can be reconstructed by means of $\text{CRSS.Combine}(\cdot)$. Otherwise, colluding users must not be able to retrieve the secret.

Instantiation. Our CRSS scheme is based on the threshold secret-sharing scheme proposed in [15], which is defined as follows:

$g^x, \{x_1, \dots, x_n\} \leftarrow \text{SS.Share}(-, t, n)$. Pick a cyclic group G of prime order q where the discrete logarithm assumption holds; let $\langle g \rangle = G$. Pick a random $x \in Z_q$ and set the secret to g^x . Pick a random $t - 1$ -degree polynomial X with coefficients in Z_q , such that $X(0) = x$. Set the i -th share to $x_i = X(i)$.

$s' \leftarrow \text{SS.Combine}(\{x_{i_1}, \dots, x_{i_l}\})$. Given shares $\{x_{i_1}, \dots, x_{i_l}\}$, use polynomial interpolation to recover the secret. That is $s' = g^{\sum_{p=1}^{p=l} x_{i_p} \lambda_p}$ where $\lambda_p = \prod_{1 \leq k \leq l, k \neq p} \frac{x_{i_k}}{x_{i_k} - x_{i_p}}$.

Note that in the above scheme, the secret is not given as input to the Share algorithm; rather, it is set to g^x for a randomly chosen x . Given the above algorithms, our CRSS scheme unfolds as follows:

$\{s, s_1, \dots, s_n\} \leftarrow \text{CRSS.Share}(-, t, n)$. Run $\text{SS.Share}(-, t, n)$ to obtain $g^x, \{x_1, \dots, x_n\}$. Pick $H(\cdot) : \{0, 1\}^* \rightarrow G$ to be a cryptographic hash function that maps random strings in G . Pick a random $t - 1$ -degree polynomial Y with coefficients in Z_q , such that $Y(0) = 0$, and denote $y_i = Y(i)$. The secret is set to $s = g^x$ while each share is set to $s_i = (x_i, y_i)$.

$d_{i,j} \leftarrow \text{CRSS.Delegate}(s_i, U_j)$. Parse $s_i = (x_i, y_i)$ and output $d_{i,j} = g^{x_i H(U_j)^{y_i}}$.

$s' \leftarrow \text{CRSS.Combine}(\{d_{i_1,j_1}, \dots, d_{i_l,j_l}\})$. Run $s' \leftarrow \text{SS.Combine}(\{d_{i_1,j_1}, \dots, d_{i_l,j_l}\})$.

Correctness. If $l \geq t$, then $\text{CRSS.Combine}(\{d_{i_1,j}, \dots, d_{i_l,j}\})$ outputs

$$\begin{aligned} s' &= \prod_{p=1}^{p=l} (d_{i_p,j_p})^{\lambda_{i_p}} = \prod_{p=1}^{p=l} (g^{x_{i_p} H(U_j)^{y_{i_p}}})^{\lambda_{i_p}} = \\ &= g^{\sum_{p=1}^{p=l} \lambda_{i_p} x_{i_p}} H(U_j)^{\sum_{p=1}^{p=l} \lambda_{i_p} y_{i_p}} = \\ &= g^k H(U_j)^0 = g^k = s. \end{aligned}$$

Security. The security of CRSS is based on the fact that, in the random oracle model, delegations for different identities cannot be combined to remove the blinding factor from the secret. Assuming that $H(\cdot)$ is modeled as a random oracle and that the discrete logarithm assumption holds in G , we can show that any p.p.t. algorithm \mathcal{A} has only negligible advantage in winning **Game Priv***.

3.5 Commune: Protocol Specification

Recall that **Commune** leverages a shared repository, which is an abstraction of the owners' storage space on \mathcal{S} . The shared repository uses a versioning system so that content cannot be overwritten but only new content can be added. In particular, **Commune** optimizes performance by splitting a file in smaller *units*, and encoding/decoding each unit separately. Therefore, when a new file version is written to the shared repository, the writer only needs to upload the units that have changed from the previous version.

Files written to the repository are encoded in *tokens* and distributed across the owners' accounts. Leveraging the basic ACLs of \mathcal{S} , owners exert their individual policy on the tokens they store on their accounts. The distributed enforcement of the **SOM** policy is implied by the enforcement of each owner's individual policy on his tokens by \mathcal{S} .

Encoding must guarantee both correctness and security of reading operations. Hence, users who are authorized to read at least t tokens must be able to decode the original file; users who are granted read access on fewer than t tokens must not be able to recover its content. Furthermore, users must not be able to pool their credentials to escalate their access rights.

Create a File. File creation requires one user, the file creator, to “bootstrap” the system and write the initial version of the file into the repository. For this reason, we assume that—at the file creation time—the file creator has been granted the right to write new data to each of the owner's accounts on \mathcal{S} .

The file creator first divides the file F into k fixed-sized units. For each unit F_i ($i \in [1, \dots, k]$), he runs $\{s_i, s_{i1}, \dots, s_{in}\} \leftarrow \text{CRSS.Share}(-, t, n)$ to produce a fresh secret s_i and n of its shares. Secret s_i is used as a symmetric key to encode the unit F_i in n chunks using SFD. That is, the file creator runs $\{c_{i1}, \dots, c_{in}\} \leftarrow \text{SFD.Encode}(t, n, F_i, s_i, \lambda)$. The token of the unit F_i for the owner O_j is set to (c_{ij}, s_{ij}) (i.e., one chunk outputted by $\text{SFD.Encode}(\cdot)$ and one secret-share outputted by $\text{CRSS.Share}(\cdot)$). Finally, for each owner O_j , the file creator writes $\{(c_{ij}, s_{ij})\}_{i \in [1, \dots, k]}$ to O_j 's account on \mathcal{S} . Each owner, therefore, receives one token for each unit that constitutes F .

Grant/Deny Write Rights. An owner O_j grants write rights to a user U_l by granting to U_l the right to write new data (i.e., tokens) to O_j 's account. Similarly, O_j denies write rights to U_l by denying U_l the right to write new data to O_j 's account.

Update a File. Assume U_l wants to write a new version of a file F . For simplicity, assume that the new version differs from the previous one by only one unit F_i (the case where the old and the new versions differ in several units is handled in a similar fashion). At this point, some owners may allow U_l to write tokens to their accounts while others may not.

Let \mathcal{O}^+ be the subset of owners who grant to U_l write rights to their accounts. Similarly, let \mathcal{O}^- be the subset of owners who deny to U_l write rights to their accounts. U_l can, therefore, only distribute tokens to owners in \mathcal{O}^+ . This scenario is equivalent to the case where U_l distributes tokens to all owners in \mathcal{O} , but the ones in \mathcal{O}^- decide to reject the version produced by U_l and make the received tokens unavailable.

U_l is an authorized writer and his version accepted (i.e., considered as written to the shared repository) if and only if $|\mathcal{O}^+| \geq t$. In this case, there are at least t tokens for the new unit, so it may be decoded by users who collect enough credentials. If $|\mathcal{O}^+| < t$, user U_l is not authorized to write and his version is rejected (i.e., considered as not written to the repository), since there are not enough tokens to decode the unit produced by U_l .

Grant/Deny Read Rights. Recall that for each unit F_i , an owner O_j receives the token (c_{ij}, s_{ij}) . O_j can grant to U_l read access to that unit by *endorsing* the token for U_l and granting to U_l read access on the endorsed token. Token endorsement requires O_j to run $d_{ij,l} \leftarrow \text{CRSS.Delegate}(s_{ij}, U_l)$. The endorsed token $(c_{ij}, d_{ij,l})$ is then made available by O_j for U_l to read. If a file consists of multiple units, O_j must endorse all relative tokens for U_l and grant to U_l read access on all endorsed tokens.

O_j can revoke read rights that were previously granted, by denying to U_l the right to read the previously endorsed tokens.

Read a File. If the original file spans several units, U_l must decode each unit separately in order to read the entire file. That is, for each unit, he uses the set of endorsed tokens he can fetch to recover the secret key via $\text{CRSS.Combine}(\cdot)$ and then uses the secret key to decode the unit via $\text{SFD.Decode}(\cdot)$. Note that for an authorized reader to read version x of file F , he must fetch the latest endorsed tokens created up to (and including) version x , for each unit that comprises the file. Assume user U_l is granted read access to $\{(c_{ij_1}, d_{ij_1,l}), \dots, (c_{ij_t}, d_{ij_t,l})\}$. To recover F_i that user runs $s_i \leftarrow \text{CRSS.Combine}(\{d_{ij_1,l}, \dots, d_{ij_t,l}\})$ and then $F_i \leftarrow \text{SFD.Decode}(s_i, \{c_{ij_1}, \dots, c_{ij_t}\}, \lambda)$. U_j proceeds in a similar way in order to recover all units of F that he has access to.

Security analysis. From Sections 3.3 and 3.4, it follows that given t tokens of a file unit F_i (endorsed for a unique user identity), it is possible to recover both the secret key used to encode F_i and its AON-FFT ciphertext, so that the original file can be decrypted. That is, users can read files written by honest writers, if they are granted such right by at least t out of n owners.

Property P1 (cf. Section 3.2) is fulfilled as follows. First, **Commune** uses copy-on-write to prevent writers from overwriting content in the shared repository with garbage. Second, malicious writers (i.e., writers who have been granted write access by fewer than t owners) are unable to distribute a file without honest readers detecting it. In other words, a file is considered as written if and only if it is correctly encoded in tokens and those tokens are distributed to and endorsed by at least t out of n owners. Any content distributed through other means (e.g., out of band channels) is easily recognized as malicious by honest readers. We argue that detection of unauthorized files is the only solution for

Parameter	Default Value
w	128 B
t	4
n	10
$ F_i $	10 MiB

Table 2: Default parameter used during the evaluation.

protecting honest readers, because there are no mechanisms to deter malicious writers from disseminating arbitrary content (cf. Observation 3). We also stress that honest readers can easily detect writers that distribute polluted (i.e., non-decodable) tokens. Denial-of-service attacks are, nevertheless, out of the scope of this work.

Property P2 is satisfied by combining CRSS and SFD. The former ensures that coalitions of users, where no single user has enough tokens endorsed for his identity, cannot pool their endorsed tokens in order to escalate their access rights. The latter addresses the case where at a time τ_1 a user has access to t or more endorsed tokens of a file unit F_i , but at a time $\tau_2 > \tau_1$, his access rights are revoked. That is, at the time τ_2 , the user has access to fewer than t endorsed tokens. SFD ensures that even if, at the time τ_1 the user may have cached the key used to encode F_i in tokens, he will not be able at the time τ_2 to decode parts of F_i . Note that, once a user has access to the file, then he can locally store any plaintext content of his choice. Similar to other access control schemes, **Commune** cannot deter this behavior.

Finally, given the guarantees that **Commune** makes for enforcing the write and read actions, it follows that **Commune** is a (correct) solution for distributed enforcement of the SOM access control policy (see Definition 1).

4. PROTOTYPE DESIGN & EVALUATION

We implement a prototype of **Commune** integrated with Amazon S3 [1]. In this section, we describe the implementation and evaluate its performance.

4.1 Implementation Setup

We leverage Amazon S3 to instantiate \mathcal{S} : for each user in \mathcal{U} , we create personal accounts in Amazon S3, into which users can upload content and define arbitrary access control policies. In our implementation, we use Amazon S3 access control features to distribute tokens from the file creator to the set of owners $\mathcal{O} \subseteq \mathcal{U}$. In particular, we assume that each user sets up (i) one “temporary” folder where other peers are granted write access, and (ii) one “main” folder where endorsed tokens are stored and retrieved. When the file creator wants to distribute a token to owner O_j , he writes the token to O_j ’s temporary folder. Since no other user apart from O_j has read access to the temporary folder, the new token is protected from unauthorized access. At this point, O_j can endorse the token for another user U_l , store the endorsed token in his main folder, and grant read access rights on the endorsed token to U_l .

Our prototype, implemented in Java, is a multi-threaded client-side interface to repositories hosted on Amazon S3. The client runs on a user’s machine and uploads/downloads content to/from the repositories. For the evaluation, we use an Intel Core i5-2400 (at 3.10 GHz), where up to 4 GB of RAM are allocated to the OpenJDK VM.

	Peak Throughput (Mbps)
Write	43.39
Read	29.52

Table 3: Peak throughput. We assume the default parameters of Table 2. Each data point is the average of 20 measurements.

The implementation of SFD leverages Rijndael [17] (implemented using the Bouncy Castle Java library [4]) as the underlying block cipher for AON-FFT and systematic Reed-Solomon codes [30] (implemented using the Jerasure library [2, 23]) for information dispersal. We chose a symbol size of 16 bytes, and a security parameter $\lambda = 128$ bits. Our implementation of CRSS leverages the secretsharejava library [3] with a 386-bit modulus.

To optimize performance, our prototype handles file unit operations at a smaller granularity, called *pieces*. During the creation of any file unit, the unit is split into pieces that are processed in parallel. A token for each unit contains one output chunk of SFD for each piece that composes the unit. The piece size w is chosen such that $t\lambda|w$, where λ is the security parameter and t is the required reconstruction threshold. This condition ensures that (i) a piece can be encrypted in an integer number of ciphertext blocks of λ bits, (ii) an encrypted piece can be divided into an integer number of input chunks for the Reed-Solomon encoder, and (iii) the size of each chunk of the Reed-Solomon encoder/decoder is at least λ bits.

4.2 Evaluating Single Unit Write/Read

We evaluate the performance of *Commune* for a single file unit write and read, with respect to (i) the piece size w , (ii) the reconstruction threshold t , (iii) the number of owners n , and (iv) the size of the file unit $|F_i|$. We assume the default parameter values shown in Table 2.⁴ We then change one variable at a time, to assess its impact on the system performance. For each configuration, we measure the time required (i) to create and upload F_i (denoted by *Unit Write* in our plots), and (ii) to retrieve F_i (denoted by *Unit Read*). These times are measured from the initiation of the operation until the output is available either in the repositories (for unit write) or on a local disk (for unit read). We control for the effect of caching by uploading random binary streams at each run.

During *Unit Read*, our client fetches endorsed tokens from t randomly chosen owners. Recall that a (t, n) systematic erasure code outputs t data chunks and $n - t$ parity chunks. Since data chunks need not be decoded, our evaluation accounts for the average-case scenario where the probability that a token contains a data chunk is bounded by $\frac{t}{n}$. Note that we do not evaluate the time required to grant read rights (i.e., the time required to endorse a token), since it does not depend on any of the considered parameters.

Our results are depicted in Figure 2. Each data point is averaged over 20 runs; where appropriate, we also provide the corresponding 95% confidence intervals. We also monitor the runtime of the intermediate steps for a number of configurations (Figure 3).

⁴Since our SFD scheme requires that the number of plaintext blocks input to AON-FFT is a power of 2, we also set w , λ , and t to be powers of 2, in order to ensure that $t\lambda|w$.

Our evaluation shows that writing a new unit (*Unit Write*) is less expensive than reading it (*Unit Read*). This is due to the overhead of thread synchronization on a single file descriptor when storing decoded pieces on the local disk.

Impact of the Piece Size: Figure 2(a) shows the impact of the piece size w on the performance of *Commune*. Smaller w leads to a smaller number of input blocks to the AON-FFT scheme, which results in better performance. Recall that AON-FFT requires $\log_2 m$ rounds of encryption over all the m input blocks. However, we experience higher latencies for very small values of w , especially in the *Unit Read* operation. This is due to the overhead incurred by different threads that must synchronize on the single file descriptor in order to write data to disk. Throughout the rest of the evaluation, we set $w = 128$ bytes, since it offers a good performance trade-off for both file read and write operations.

Impact of the Reconstruction Threshold: Figure 2(b) shows the impact of the reconstruction threshold t on the system’s performance. As t decreases, the chunk size of the Reed-Solomon encoder increases; this results in larger chunk upload and download times. Figure 3 also shows that a smaller value of t results in longer encoding and decoding times. On the one hand, during *Unit Write*, small values of t result in larger encoding overhead since the size of the encoding matrix increases. On the other hand, during *Unit Read*, small values of t decrease the probability of recovering data chunks (w.r.t. the probability of recovering parity chunks), which makes decoding slower (cf. Figure 3).


Impact of the Number of Owners: Figure 2(c) shows that latency increases for both *Unit Write* and *Unit Read* as the number of owners grows. During *Unit Write*, this increase is caused by the distribution of tokens from the file creator to the set of owners. Latency increase during *Unit Read* is due to an higher probability of fetching parity codes that take more time to be decoded by the Reed-Solomon decoder.

Impact of the Unit Size: Figure 2(d) shows *Commune*’s performance for different unit sizes. In particular, we vary the number of pieces, that comprise a file unit. Our results show that the time required to write/read a unit increases almost linearly with the unit size. (Note that Figure 2(d) uses semi-logarithmic axes.) The time required to read a 10 MB unit is roughly 4.47 seconds. In this case, the effective throughput of our prototype is close to 18 Mbps.


To optimize the performance, the choice of the unit size should depend on the user update patterns. Larger unit sizes mean that users have to upload larger amounts of data when updating any part of the file. Smaller unit sizes result in smaller upload times for small updates. Small unit sizes, however, incur considerable overhead when a user updates significant parts of the file (i.e., when the update affects a large number of units). In the following, we study the performance of writing and reading multiple units.

4.3 Multiple units


We now assess the performance of writing/reading multiple units of a file. In particular, we measure the peak throughput exhibited by our prototype implementation. We increase the number of units that are concurrently written/read to/from Amazon S3 until the throughput is sat-




(a) Impact of the piece size.



(b) Impact of the reconstruction threshold.



(c) Impact of the number of owners.



(d) Impact of the unit size.

Figure 2: Performance evaluation of our prototype implementation. The system parameters are chosen from Table 2. Each data point is averaged over 20 measurements; where appropriate, we also provide the corresponding 95% confidence intervals.

urated. We then compute the peak throughput as the maximum aggregated amount of data (in bits) that can be written/read per second to/from Amazon S3. Table 3 shows that the peak throughput is above 29 Mbps for both write and read operations.

In summary, we conclude that **Commune**’s overhead can be tolerated in applications scenarios like authoring applications, where users work on content on their local machines and periodically upload/download content to/from the cloud. Furthermore, in a practical use case, users bear the full cost of file write/read only once. Thanks versioning, once a user has uploaded/downloaded the entire file, subsequent updates can be made to individual units.

5. RELATED WORK

Secret Sharing and Information Dispersal: Secret sharing schemes [10] allow a dealer to distribute a secret among a number of shareholders, such that only authorized subsets of shareholders can reconstruct the secret. In threshold secret sharing schemes [15, 28], the dealer defines a threshold t and each set of shareholders of cardinality equal to or greater than t is authorized to reconstruct the secret. Se-

cret sharing guarantees security (i.e., the secret cannot be recovered) against a non-authorized subset of shareholders; however, they incur a high computation/storage cost, which makes them impractical for sharing large files.

Rabin [24] proposed an information dispersal algorithm with smaller overhead than that of [28], however, his proposal does not provide any security guarantees when a small number of shares (fewer than the threshold) are available. Krawczyk [19] combines both Shamir’s [28] and Rabin’s [24] approaches; in [19] a file is first encrypted using AES and then dispersed using the scheme in [24], while the encryption key is shared using the scheme in [28].

Information dispersal based on erasure codes [30] are effective tools to enhance the reliability of cloud-based storage systems [7, 8, 20, 31]. Ramp schemes [11] constitute a trade-off between the security guarantees of secret sharing and the efficiency of information dispersal algorithms.

All or Nothing Transformations: All-or-nothing transformations were first introduced in [26] and later investigated in [13, 16]. The majority of AONTs leverage a secret key that is embedded in the output blocks. Once all output blocks are available, the key can be recovered and single




Figure 3: Runtime analysis for five different configurations of our prototype implementation. The label “Writeback” denotes the operation of writing the data back on disk onto the file descriptor (during Unit Read).

blocks can be reverted. Rivest [26] also mentioned a transformation that is inspired by Fast Fourier Transform. Van Dijk et al. [29] later on leveraged Rivest’s transformation to construct a “proof of encryption” of files in the cloud. In this paper, we extend the use of Rivest’s transformation to construct an AONT scheme, that keeps its all-or-nothing property even if the adversary is given the secret key. Resch et al. [25] combine AONT and information dispersal to provide both fault-tolerance (i.e., decoding requires only t out of n shares) and data secrecy (i.e., confidentiality is guaranteed w.r.t. parties that collect fewer than t shares), in the context of distributed storage systems. In [25], however, an adversary who caches the encryption key can still decode single shares. In [18], Karame *et al.* showed that by first encrypting the data then post-processing it using a linear transform, one can construct an *encryption mode* which provides similar guarantees as all or nothing transforms, and with comparable performance.

Access Control Systems: Current state-of-the-art access control systems, such as SecPAL [9], KeyNote [12], and Delegation Logic [22], can in principle express t out of n policies. These languages, however, rely on the presence of a centralized PDP component to evaluate their policies. Furthermore, their PDPs cannot be deployed within a third-party cloud platform. As explained in Section 2, these access control systems rely on an administrator to define and manage access control policies. In our setting, this means that a set of owners has to elect one enforcer who has unilateral powers over their files.

6. CONCLUSION

Even though existing clouds platforms are used as shared repositories, they surprisingly do not support any notion of shared ownership. We consider this to be a severe limitation because parties which contribute data to the repository cannot jointly decide how their resources are used. The problem of enforcing shared ownership in the cloud is even more difficult since a cloud platform does not allow deployment of a third-party enforcement component.

In this paper, we introduced a novel concept of shared ownership and we described it through a formal access control model, called SOM. We also proposed our scheme, Com-

mune, that distributively enforces SOM. Commune can be used in existing clouds without any modifications to the platforms. We implemented and evaluated the performance of our solution within Amazon S3. Our results show that Commune scales well with the file size and with the number of users.

Given the rise of personal clouds (e.g., [5, 6]), we argue that Commune finds direct applicability in setting up shared repositories that are distributively managed atop of the various personal clouds owned by users. We therefore hope that our findings motivate further research in this area.

7. REFERENCES

- [1] Amazon Simple Storage Service (Amazon S3). <http://aws.amazon.com/s3/>.
- [2] Github—Jerasure. <https://github.com/tsuraan/Jerasure>.
- [3] Shamir’s Secret Share in Java (secretsharejava). <http://sourceforge.net/apps/trac/secretsharejava/wiki>.
- [4] The Legion of the Bouncy Castle. <http://www.bouncycastle.org/java.html>.
- [5] The Respect Network. <https://www.respectnetwork.com/>.
- [6] WD My Cloud. <http://www.wdc.com/en/products/products.aspx?id=1140>.
- [7] M. Abd-El-Malek, G. R. Ganger, G. R. Goodson, M. K. Reiter, and J. J. Wylie. Fault-Scalable Byzantine Fault-Tolerant Services. In *ACM Symposium on Operating Systems Principles (SOSP)*, pages 59–74, 2005.
- [8] M. K. Aguilera, R. Janakiraman, and L. Xu. Using Erasure Codes Efficiently for Storage in a Distributed System. In *International Conference on Dependable Systems and Networks (DSN)*, pages 336–345, 2005.
- [9] M. Y. Becker, C. Fournet, and A. D. Gordon. SecPAL: Design and Semantics of a Decentralized Authorization Language. In *Journal of Computer Security (JCS)*, pages 597–643, 2010.
- [10] A. Beimel. Secret-sharing schemes: A survey. In *Third International Workshop on Coding and Cryptology (IWCC)*, pages 11–46, 2011.

- [11] G. R. Blakley and C. Meadows. Security of ramp schemes. In *Advances in Cryptology (CRYPTO)*, pages 242–268, 1984.
- [12] M. Blaze, J. Ioannidis, and A. D. Keromytis. Trust Management for IPsec. In *ACM Transactions on Information and System Security (TISSEC)*, pages 95 – 118, 2002.
- [13] V. Boyko. On the Security Properties of OAEP as an All-or-nothing Transform. In *Proceedings of CRYPTO*, pages 503–518, 1999.
- [14] S. Ceri, G. Gottlob, and L. Tanca. What you always wanted to know about Datalog (and never dared to ask). In *Knowledge and Data Engineering, IEEE Transactions on*, pages 146 –166, 1989.
- [15] C. Charnes, J. Pieprzyk, and R. Safavi-Naini. Conditionally secure secret sharing schemes with disenrollment capability. In *ACM Conference on Computer and Communications Security (CCS)*, pages 89–95, 1994.
- [16] A. Desai. The security of all-or-nothing encryption: Protecting against exhaustive key search. In *Advances in Cryptology (CRYPTO)*, pages 359–375, 2000.
- [17] J. Daemen, and V. Rijmen. AES Proposal: Rijndael. <http://csrc.nist.gov/archive/aes/rijndael/Rijndael-ammended.pdf>.
- [18] G. O. Karame, C. Soriente, K. Lichota, and S. Capkun. Securing cloud data in the new attacker model. *IACR Cryptology ePrint Archive*, 2014:556, 2014.
- [19] H. Krawczyk. Secret Sharing Made Short. In *International Conference on Advances in Cryptology*, 1993.
- [20] J. Kubiawicz, D. Bindel, Y. Chen, S. E. Czerwinski, P. R. Eaton, D. Geels, R. Gummadi, S. C. Rhea, H. Weatherspoon, W. Weimer, C. Wells, and B. Y. Zhao. OceanStore: An Architecture for Global-Scale Persistent Storage. In *International Conference on Architectural Support for Programming Languages and Operating Systems (ASPLOS)*, pages 190–201, 2000.
- [21] A. B. Lewko and B. Waters. Decentralizing Attribute-Based Encryption. In *International Conference on the Theory and Application of Cryptographic Techniques (EUROCRYPT)*, pages 568–588, 2011.
- [22] N. Li, B. N. Grosz, and J. Feigenbaum. Delegation logic: A Logic-based Approach to Distributed Authorization. In *ACM Transactions on Information and System Security (TISSEC)*, pages 128–171, 2003.
- [23] J. S. Plank, S. Simmerman, and C. D. Schuman. Jerasure: A library in C/C++ facilitating erasure coding for storage applications. Technical report, 2007.
- [24] M. O. Rabin. Efficient Dispersal of Information for Security, Load Balancing, and Fault Tolerance. In *Journal of the Association for Computing Machinery*, pages 335–348, 1989.
- [25] J. K. Resch and J. S. Plank. AONT-RS: Blending Security and Performance in Dispersed Storage Systems. In *USENIX Conference on File and Storage Technologies (FAST)*, pages 191–202, 2011.
- [26] R. L. Rivest. All-or-Nothing Encryption and the Package Transform. In *International Workshop on Fast Software Encryption (FSE)*, pages 210–218, 1997.
- [27] P. Rogaway and M. Bellare. Robust computational secret sharing and a unified account of classical secret-sharing goals. In *ACM Conference on Computer and Communications Security (CCS)*, pages 172–184, 2007.
- [28] A. Shamir. How to Share a Secret? In *Communications of the ACM*, pages 612–613, 1979.
- [29] M. van Dijk, A. Juels, A. Oprea, R. L. Rivest, E. Stefanov, and N. Triandopoulos. Hourglass Schemes: how to prove that cloud files are encrypted. In *ACM Conference on Computer and Communications Security (CCS)*, pages 265–280, 2012.
- [30] J. H. van Lint. *Introduction to Coding Theory*. Springer-Verlag New York, Inc., Secaucus, NJ, USA, 1982.
- [31] H. Xia and A. A. Chien. RobuStore: a Distributed Storage Architecture with Robust and High Performance. In *ACM/IEEE Conference on High Performance Networking and Computing (SC)*, page 44, 2007.

APPENDIX

A. DATALOG ENCODING OF THE SOM ACCESS CONTROL POLICY

To show that SOM can be centrally enforced, we present its encoding in the Datalog logic-programming language.

We represent a file system state as a Datalog database [14] that has a set of relations describing each file’s owners and its threshold. We translate requests and credentials into Datalog clauses (rules), which are then evaluated over the current state together with the SOM policy. We first give a brief overview of Datalog (see [14] for a more extensive survey). A Datalog program is a finite set of clauses of the form: $S \leftarrow L_1, \dots, L_n$, where S and L_i are function-free first-order literals of the form $predicate(arg_1, \dots, arg_n)$. We refer to S as the head of the clause, and to L_i as a body literal. We adopt the following notation: a variable starts with the ? character, a constant starts with a capital letter, and a predicate name starts with a lower-case letter. A clause with no body literals is called a *fact*. All clauses are safe: all variables that appear in a head literal also appear in at least one body literal. A Datalog program can be split into two sets of clauses: *EDB* and *IDB*. *EDB* is a set of facts whose head literals do not appear as head literals in any other clause. All other clauses are in the *IDB* set. Intuitively, we think of an *EDB* as an input for computing all implied facts by the clauses in the *IDB* set. The declarative semantics of a Datalog program interpret each clause as a first-order sentence: $\forall \bar{x} L_1 \wedge \dots \wedge L_i \rightarrow S$, and take a whole program to be a conjunction of its clauses. For each program $\mathcal{P} = IDB \cup EDB$, let $\sigma(IDB, EDB) = \{atom \mid \mathcal{I}(\mathcal{P}) \models atom\}$, where $\mathcal{I}(\mathcal{P})$ is the first-order translation of \mathcal{P} , and \models is the logical implication.

A SOM state s is a tuple (*Files*, *Users*, *Owens*, *Thresholds*) where *Files* denotes a set of strings representing file names, *Users* is a set of users, *Owens* is a subset of $2^{Users \times Files}$, and *Thresholds* is a mapping from *Files* into \mathbb{N} . For a state s , we define a set EDB_s containing all ground atoms: $file(File)$, $user(User)$, $owns(User, File)$, and $threshold(File, N)$. A request r is a tuple (R, \mathcal{C}), where R is a request credential submitted by a user, and \mathcal{C} is a set of available credentials.

Credentials can be either submitted by a user, or kept in a separate storage and appended to each request. Given a request, $\mathcal{T}(C)$ generates the following set of Datalog rules IDB_r :

$$\mathcal{T}(O \text{ says } U \text{ can } actionOp(F)) = says(O, U, actionOp, F)$$

The translation of R is similar, except that we do not generate *says* facts but Datalog queries:

$$\mathcal{T}(U \text{ reqs } actionOp(F) = can(U, actionOp, F)$$

The set \mathcal{A}_s is a set of Datalog rules *parameterized* on s , that enforces the shared ownership:

$$\begin{aligned} can(?U, actionOp, ?F) \leftarrow file(?F), user(?U), \\ threshold(?F, ?T), \\ [[says(?O_1, ?U, actionOp, ?F), \dots, \\ says(?O_T, ?U, actionOp, F), \\ owns(?O_1, ?F), \dots, owns(?O_T, ?F), \\ ?O_1 \neq ?O_2, \dots, ?O_1 \neq ?O_T, \dots, ?O_{T-1} \neq ?O_T]] \end{aligned}$$

Intuitively, the given rule is a *template* rule that instantiates the necessary clauses for all *actionOp* operations. The variable $?T$ denotes a threshold, and $?U$ denotes a user. The reason for doing so is to correctly enforce the current (for the given state s) threshold t for a particular file. In short, we need to generate the correct number of $?O_i$ variables for each file and its threshold in s . To represent this *dynamic* part of a clause, we enclose it within $[[$ and $]]$ brackets.

Finally, given a SOM state s , and a request (R, C) , we say that SOM's Datalog-based Policy Decision Point (PDP) grants R if and only if: $\mathcal{I}(EDB_s \cup \mathcal{A}_s \cup \mathcal{T}(C)) \models \mathcal{T}(R)$

B. SECURITY ANALYSIS OF SFD

We treat the information dispersal encoder of SFD.Encode(\cdot) as a ramp scheme [11]. In particular, to maximize the code rate, we assume a $(0, t, n)$ -ramp scheme, with t as reconstruction threshold and 0 as privacy threshold. This means that any single chunk leaks information about the encoded input. Let \bar{F} denote the input to the information dispersal encoder. Hence, at least t chunks are necessary to reconstruct \bar{F} ; however, each single chunk leaks some information about \bar{F} . More specifically, if the adversary \mathcal{A} is given l out of n chunks outputted by SFD.Encode(t, n, F, K, λ), then:

$$H(\bar{F}|\mathcal{A}) = \begin{cases} \frac{l}{t} H(\bar{F}) & \text{if } l < t \\ 0 & \text{if } l \geq t \end{cases}$$

where $H(\cdot)$ denotes the entropy. Since \bar{F} is the output of AONT (which is essentially a block cipher) we can assume $H(\bar{F}) = m\lambda$. Therefore, if \mathcal{A} is given $l = t - 1$ chunks, $H(\bar{F}|\mathcal{A}) = \frac{m}{t}\lambda$. For simplicity, we assume that when $m = t$ and given $t - 1$ SFD encoded chunks, \mathcal{A} can decode all but one block of the AONT output.

We are left to show that, given K and $\bar{f}_{i_1}, \dots, \bar{f}_{i_{m-1}}$, \mathcal{A} has negligible advantage in recovering $f \subseteq F$, i.e., a substring of F of size λ . We note that the security of the underlying block-cipher prevents the adversary from recovering partial bits of any cleartext block. That is, the adversary can only learn entire blocks of cleartext. Therefore, we focus our analysis on an adversary that tries to recover any cleartext block $f_i \in F$. Furthermore, in our definition the adversary is not allowed to store any blocks of ciphertext/cleartext nor

is she allowed to store any intermediate block produced by AON-FFT. We argue that if an adversary were to store any block of data, then she would store blocks the actual file to easily access it despite revocation. No mechanism can cater for effective revocation if the adversary has a local copy of the resource.

We prove the security of AON-FFT by induction. Denote with AON-FFT(4) the graph that defines the operations of AON-FFT when $m = 4$. Clearly, recovering any input block of AON-FFT(4) requires all 4 output blocks. Recall that each output block is λ -bit in size and has high-entropy (since it is the output of a block-cipher). Therefore, given all but one output block, \mathcal{A} has only negligible advantage in guessing the missing block and recovering any input block. Similarly, an AON-FFT(8) graph has two AON-FFT(4) subgraphs (one left and one right), plus an additional round of encryption. Recovering any input block of AON-FFT(8) requires one input blocks of each of the two AON-FFT(4) subgraphs which, in turn, require all their output blocks. Therefore, recovering any input block of AON-FFT(8) requires all 8 output blocks. By iterating this analysis for larger graph sizes, we can easily prove that for any m , recovering any input block of AON-FFT(2^m) requires all 2^m output blocks.

C. SECURITY ANALYSIS OF CRSS

We now prove the security of our CRSS scheme (Section 3.4). More specifically, we show that a p.p.t. algorithm \mathcal{A} that has non-negligible advantage in winning **Game Priv*** of Section 3.4, can be used by p.p.t. algorithm \mathcal{B} as an internal routine to break the security of the threshold secret sharing scheme in [15]. That is, \mathcal{B} is challenged by a challenger C to find the secret g^x given only $t - 1$ shares as output by $g^x, \{x_1, \dots, x_n\} \leftarrow \text{SS.Share}(-, t, n)$. On the other hand, \mathcal{B} challenges \mathcal{A} to break the security of CRSS. Simulation starts with \mathcal{A} who submits two messages x_0, x_1 of equal length. Those are forwarded by \mathcal{B} to C that flips an unbiased coin b and secret-shares x_b . Given our particular secret sharing scheme, we assume that the secret to be reconstructed is g^{x_b} . At this time, \mathcal{B} also picks a random polynomial Y of degree $t - 1$, such that $Y(0) = 0$ and computes $y_i = Y(i)$, for $1 \leq i \leq n$.

During Type-1 queries, \mathcal{A} submits indexes $i_1, \dots, i_{t'}$. \mathcal{B} forwards them to C that replies with $\{x_{i_1}, \dots, x_{i_{t'}}\}$. \mathcal{B} then sends to \mathcal{A} shares $\{s_{i_1}, \dots, s_{i_{t'}}\}$ where $s_{i_l} = (x_{i_l}, y_{i_l})$, for $1 \leq l \leq t'$.

Similarly, during a Type-2 query \mathcal{A} submits a fresh identity U_j and t'' indexes $i'_1, \dots, i'_{t''}$. For each identity U_j , \mathcal{B} picks at random $h_j \in Z_q$ and sets $H(U_j) = g^{h_j}$. If an index i'_l has been submitted during type-1 query, \mathcal{B} knows $x_{i'_l}$ and can compute the delegation $d_{i'_l, j} = g_{i'_l}^{x_{i'_l}} H(U_j)^{y_{i'_l}}$. Otherwise, \mathcal{B} asks C for share $x_{i'_l}$ and computes the corresponding delegation.

Note that \mathcal{A} can submit a Type-2 query if only if $t' < t - 1$. Therefore \mathcal{B} can still ask C for the missing shares. During the guess stage, \mathcal{A} will output his guess b' and \mathcal{B} will use it as its own guess towards C . Since $H(\cdot)$ is a random oracle and \mathcal{A} has a non-negligible advantage in guessing b , then \mathcal{B} has the same advantage in breaking the security of [15], thus concluding our proof.