

RITM: Revocation in the Middle

Pawel Szalachowski, Laurent Chuat, Taeho Lee, and Adrian Perrig
Department of Computer Science
ETH Zurich, Switzerland

Abstract—Although TLS is used on a daily basis by many critical applications, the public-key infrastructure that it relies on still lacks an adequate revocation mechanism. An ideal revocation mechanism should be inexpensive, efficient, secure, and privacy-preserving. Moreover, rising trends in pervasive encryption pose new scalability challenges that a modern revocation system should address. In this paper, we investigate how network nodes can deliver certificate-validity information to clients. We present RITM, a framework in which middleboxes (as opposed to clients, servers, or certification authorities) store revocation-related data. RITM provides a secure revocation-checking mechanism that preserves user privacy. We also propose to take advantage of content-delivery networks (CDNs) and argue that they would constitute a fast and cost-effective way to disseminate revocations. Additionally, RITM keeps certification authorities accountable for the revocations that they have issued, and it minimizes overhead at clients and servers, as they have to neither store nor download any messages. We also describe feasible deployment models and present an evaluation of RITM to demonstrate its feasibility and benefits in a real-world deployment.

I. INTRODUCTION

Checking the revocation status of a certificate is a fundamental process for the security of any connection established via a public-key infrastructure (PKI). Certificate revocation occurs when a non-expired certificate must be invalidated. In order to revoke a certificate, the certification authority (CA) that issued the certificate adds its *serial number*¹ to a *blacklist*. Then, the CA tries to publish or disseminate the most recent version of that list. This process is sometimes referred to as the Grand Challenge of PKIs [22], for several reasons: 1) To keep connections secure, revocations must be accessible and the dissemination should be fast, so that every interested party is able to deny an illegitimate connection (that relies on a revoked certificate) as soon as possible. 2) The revocation system should be robust and resilient to server misconfiguration and blocking attacks. 3) The revocation status should be authentic (i.e., all clients can verify that any given revocation was created by an authorized CA). 4) The system should preserve user privacy (i.e., no third party, such as a CA, should be able to learn about the servers that clients tried to contact).

Several revocation schemes have been proposed and deployed over the last few years, but the current situation is still unsatisfactory in many aspects. For instance, CRLs and OCSP violate user privacy and are inefficient, while the deployment of OCSP stapling requires server reconfiguration. Due to these drawbacks, browser vendors recently decided to push special revocation lists to clients using software update mechanisms.

¹The serial number is a positive integer value (represented at most by 20 bytes), assigned uniquely to every CA-issued certificate [15].

We consider that such a scheme is not future-proof, as even now the revocation lists issued by these vendors are limited to a small fraction of all certificates (for efficiency reasons), and the number of certificates/revocations is expected to grow.

Revocation systems will face new scalability challenges in the near future, due to the increasing security awareness of website owners and hosting providers, and new initiatives that aim to make TLS ubiquitous. For instance, the *Let's Encrypt* [8] project provides a free certificate to any website administrator who requests it, as well as a convenient certificate-management software. The project, during few months of operation, issued more than one million certificates which protect approximately 2.4 million domains [8]. The *Universal SSL* project [32] initiated by CloudFlare constitutes another prominent example of mass TLS adoption. It provides over 2 million clients with free TLS certificates. Furthermore, hardware-limited/mobile devices and paradigms such as the Internet of Things (IoT) may rapidly increase the number of TLS clients. As a consequence, we may expect that the number of TLS connections will increase by orders of magnitude, in the near future. Unfortunately, mass TLS adoption poses scalability problems for current revocation systems, as pervasive TLS deployment would result in an increased number of revocations. As reported by Liu et al. [27], even today, due to efficiency reasons, all mobile browsers have disabled revocation checks.

Catastrophic events, such as the disclosure of the infamous Heartbleed vulnerability [46], [19], may again result in a huge number of revocations. In the current setting, revocation propagation, besides operational issues, reportedly induces significant financial costs for CAs [31]. This might lead to a situation in which mass-revocation decisions are strongly influenced by financial concerns.

In this paper, we present RITM, the first scheme to incorporate revocation functionalities into middleboxes. The main intuition behind our solution is that clients do not have to store any revocation lists nor need to establish any new connection (to check revocation status) if all of their TLS traffic goes through a middlebox that can provide them with the correct revocation status. RITM consists of two mechanisms: a dissemination network that provides revocation messages to middleboxes, and a protocol that allows middleboxes on a TLS client-server path to deliver revocation status to clients.

RITM minimizes overhead at clients and servers, and can be easily implemented with low operational costs as part of an efficient and profitable deployment model. The system is resilient against a range of well-known attacks against the

current revocation schemes (like blocking or server reconfiguration). Moreover, following the transparency trend of the PKI ecosystem, RITM keeps CAs accountable for the revocations they issue, and it minimizes the probability of a successful attack in which a CA would maliciously alter a revocation. We evaluate RITM by showing its performance, overhead, cost-effectiveness, and feasibility in a real-world deployment. We also demonstrate that RITM can disseminate a new revocation within seconds to all supported parties across the Internet.

II. BACKGROUND AND RELATED WORK

Hash Chains and Hash Trees. We denote a cryptographic hash function as $H(\cdot)$. A *hash chain* of length m is the successive application of a hash function $H(x)$ (m times) to an input x . The m -th value of a hash chain is denoted as $H^m(x)$, for example $H^3(x) = H(H(H(x)))$ and $H^0(x) = x$. Using hash chains it is possible to compute $H^m(x)$ from any $H^i(x)$ if $i < m$, but it is infeasible to compute $H^m(x)$ from any value $H^j(x)$ where $j > m$. In a *hash tree* every non-leaf node is a hash of its child nodes. This structure allows to efficiently (logarithmically in the number of leaves) prove that a given node is an element of the tree. When the leaves are sorted in a predefined order, it is also possible to efficiently prove that a given element is absent.

Content-Delivery Network. A CDN is a distributed infrastructure dedicated to efficiently deliver content to clients. CDNs are mostly driven by web content, and their main goal is to reduce latency and the load exerted on *distribution points*. To operate efficiently, a CDN consists of many *edge servers*, geographically distributed in the proximity of large collections of clients. Edge servers replicate the original content by obtaining it from a distribution point. The dominant communication paradigm used in CDNs is *data pulling*, i.e., clients pull the content from edge servers, which similarly pull it from the distribution points. To work efficiently, edge servers cache the content for a specific *time to live* (TTL), which usually is set by the origin. From the client's perspective, the CDN is an abstraction, as routes to edge servers are not made explicit.

Desired Properties. The revocation system must be able to handle *near-instantaneous revocations*, i.e., revocation information should be accessible by the interested parties within seconds. This property is crucial to minimize the *attack window*, which, with the current revocation systems, varies from hours to days. A revocation system must have the ability to meet this requirement in order to support applications such as e-banking, online trading, or transaction processing from point-of-sale terminals. We stress that this property implies that a client should also be notified about a revocation in the middle of established connections. Otherwise, a *race condition* is possible when the client establishes a connection just before the corresponding certificate is revoked. This is especially important for applications that establish long-lived TLS connections, like VPNs, TLS tunnels, remote filesystems, secure messaging, or environments like IoT (where devices are reluctant to conduct costly TLS handshakes).

The revocation process must not require *server reconfiguration/update*, or availability of an *additional third party*. Ideally, whenever a CA has initiated the dissemination of a revocation message, and a server is accessible, then the clients that connect to this server should be able to access the message. Whenever the effectiveness of a scheme depends on the configuration of the server, an adversary who compromised the server can decrease the effectiveness of the scheme by tuning configuration parameters.

Revocation information must be *authentic*. Only the legitimate party (i.e., a CA that issued a certificate) can create a fresh revocation message, and everyone must be able to verify that message. Moreover, certificate revocation status should be obtained without sacrificing *privacy*. Especially, it should not be required that clients contact a third party, as it reveals which domains they contacted.

The system must be *efficient* in terms of transmission, computation, and storage. It should not introduce any significant latency to a client-server connection. The solution must *scale* with an increasing number of revocations, certificates, servers, CAs, and clients. The revocation system must also provide or fit into *feasible and profitable deployment* models.

Following current trends in the field of PKIs, CA actions related to revocation issuance must be *transparent* and *accountable*. Anyone must be able to verify that the appropriate CA issued a given revocation, and the system must provide means to ensure that everyone has the same view of revocations. This guarantees that a whole range of potential attacks conducted by a malicious CA will be detected.

Related Work and its Drawbacks. The first attempt to address the revocation problem, proposed to distribute the information through *Certificate Revocation Lists* (CRLs) [15], published by CAs at CRL distribution points (specified in the certificate). To verify a certificate's validity, a client downloads a CRL and checks whether the certificate is listed in the CRL. Unfortunately, this approach is inefficient as an entire CRL must be downloaded to verify a single certificate, and some CRLs are currently large (a few megabytes). As this download is performed during a TLS handshake, it significantly increases the latency of the entire process. *Delta CRLs* can reduce this large overhead by allowing clients to only fetch new revocations when required, but they still suffer from other drawbacks of regular CRLs. In particular, a CA can violate user privacy by creating a dedicated distribution point for a targeted certificate, and simply determine which client contacts a targeted domain. Furthermore, Gruschka et al. report that during a 3-month period, only 86.1% of the CRL distribution points were available [21].

The *Online Certificate Status Protocol* (OCSP) [38] was proposed to address the inefficiency of CRLs. With OCSP, a client can contact a CA to obtain the current revocation status of a single, given certificate. This solution, however, is still inefficient because the CA may be under heavy load, and an extra connection during the TLS handshake is still required. Furthermore, OCSP has a serious privacy issue, as the CA learns which server the client connects to.

To solve these problems *OCSP Stapling* was proposed [30]. With this solution, the server periodically fetches an authenticated OCSP response from the CA, and then sends the response stapled along with the certificate in subsequent TLS connections. Unfortunately, this solution requires servers to update their software. Although a study suggests that some of the major domains (i.e., 22.5% of domains that are sampled from the Alexa’s list of most popular sites) have deployed OCSP stapling [41], the overall deployment rate is still marginal—a study showed that only 3% of all certificates are served by servers that adopted OCSP stapling [27]. Moreover, the age of the stapled response can be customized by a configuration parameter; therefore, a long attack window can be introduced by an adversary or a misconfiguration [42].

Short-Lived Certificates (SLCs) [36], [45] solve problems associated with CRLs and OCSP by eliminating revocation completely. SLCs are designed to be valid for a few days, which leaves a long attack window as SLCs are irrevocable for that time period. In addition, their deployment depends on software update and configuration on the server side, as every server must be set up to contact a CA periodically.

Recently, browser vendors decided to disseminate special CRLs (CRLSet [25], OneCRL [5]) through software updates and major browsers use these lists as the main revocation mechanism. Such an approach does not require any server update or reconfiguration; however, a long attack window still exists, as software updates are infrequent and as clients apply them at irregular points in time (with a heavy-tail distribution). This technique also faces scalability issues. Every client must store a complete list, which may be infeasible or undesired for hardware-limited devices [27]. As the dissemination is done through unicast communication, scalability in the face of an increasing number of clients is challenging. Finally, with an increasing number of revocations, CRLSets may become too large to be delivered to all clients. To address this issue, CRLSets include only a marginal number of revocations (0.35%, as reported [27]). Such a policy restricts the method’s scalability and effectiveness. For this reason, the CLRSet approach has been criticized [7], [27].

RevCast [39] is another recent approach, which improves revocation dissemination through unique properties of radio broadcast. RevCast proposes an architecture in which CAs broadcast revocation messages, and clients with radio receivers can receive and collect them immediately. RevCast is still a CRL-like approach where clients must possess the entire CRL. To satisfy this requirement, an additional infrastructure must be created (e.g., a device must obtain and store new CRL entries when clients are not listening to the broadcast transmission). RevCast also requires that users purchase and install radio receivers. Another downside of RevCast is that it requires to significantly change CA operations, and the maximum bandwidth is 421.8 bit/s, which prevents handling high revocation rates rapidly.

Log-based approaches [26], [24], [11], [43] make the revocation issuance process transparent and accountable, as CAs are obligated to submit revocations to public and veri-

fiable logs. Unfortunately, the deployment of these schemes is mainly server-driven (i.e., servers have to be configured to fetch and serve fresh revocation status) or client-driven (i.e., clients contact logs, which compromises their privacy). Moreover, the attack window is large, as logs are designed to update their internal state every few hours.

System Model. We investigate the revocation problem for a TLS communication in which only the server is authenticated to the client. Such a setting is typical on today’s Internet. We use TLS as our default environment, but we stress that our scheme can be combined with any protocol that exchanges certificate(s) in plaintext.² We assume that 1) CAs can reach a dedicated dissemination network at regular intervals (denoted Δ)³, 2) TLS and the cryptographic primitives that we use are secure, and 3) the different parties are loosely time synchronized. Time is expressed in Unix seconds and the *time()* function returns the current time. We denote a message *msg* signed by *X*’s private key K_X^- as $\{msg\}_{K_X^-}$.

Adversary Model. We assume that an adversary can control the network (can modify, block, and create any message) and can compromise any of its elements. The adversary’s goal is to convince non-compromised clients either that a revoked certificate is still valid or that a valid certificate is revoked.

III. RITM

Lessons Learned and Insights. We summarize some of the most important lessons learned from the deployment of previous revocation schemes (for more details see §II): 1) *During the connection establishment process a client cannot make any dedicated connection to check whether a certificate is revoked.* Otherwise, an additional latency is introduced, privacy may be violated, and the connection can be established only when the third party is available. (See CRL and OCSP.) 2) *Due to large bandwidth and storage overheads, clients cannot be constantly equipped with complete lists of revocations.* Moreover, due to the expected increase in the number of TLS clients and servers, this might become even more problematic in the future. (See CRL and CRLSet.) 3) *It is very unlikely that a significant fraction of all servers be rapidly updated to use a new security enhancement.* (See OCSP Stapling and log-based approaches.)

Our first observation is that, as clients and servers cannot be updated with fresh revocations and clients should avoid obtaining revocation status via a dedicated connection, the only remaining possibility to provide revocation status to clients is to implement a network functionality that will support it.

The second observation is that an efficient, robust, and fast dissemination network is a necessary component of a satisfactory revocation system in the current TLS ecosystem, and CDN infrastructures fulfill all these requirements.

Finally, the current formats of revocation lists do not allow to efficiently monitor the revocations issued by CAs. For

²Protocols that encrypt handshakes (like the draft version of TLS 1.3 [34]), would need to expose a server’s certificate or its identifier in plaintext.

³The value of Δ is a trade-off between the size of the attack window and efficiency. Throughout the paper we analyze values from 10 seconds to 1 day. For the sake of simplicity, we assume that Δ is a global parameter, but this assumption can be easily relaxed—see §VIII.

instance, it is challenging to detect a misbehavior when a CA shows different CRLs (or OCSP messages) to different clients (for instance to hide the fact that a given certificate is revoked). **High-level Design.** Taking into consideration the lessons learned and our observations, the high-level idea is to 1) propagate revocations through a CDN, 2) piggyback revocation status to standard TLS communications, and 3) enhance the format of revocation lists to enable a more efficient monitoring of CAs. In this setting, every revocation is distributed through a dissemination network, and for every TLS connection, a dedicated middlebox (that is connected to the dissemination network) on the path between the client and the server will put fresh, authentic, and accountable revocation status along with the native TLS message towards the client.




Fig. 1: High-level architecture of RITM.

RITM's high-level architecture is depicted in Fig. 1. The essential element of RITM is a network middlebox called the *Revocation Agent* (RA). RAs are designed to fulfill several purposes. First, they connect to the dissemination network used to propagate revocations. As opposed to client machines, middleboxes are usually always turned on, which allows them to constantly receive and collect new revocations from the dissemination network (steps 1, 2, 3). Then, for every TLS connection, an RA provides fresh revocation status to the RITM-supported clients (step 4). For non-supported clients and for non-TLS traffic, RAs act as transparent middleboxes. Finally, RAs are able to keep CAs accountable, as they monitor the revocations they issue, compare their views of revocations with other parties of the system, and report any detected misbehavior in a provable manner (step 5).

Revocation Lists. One challenge is to find a format of revocation list that meets our requirements. Ideally, it should be realized in a way that allows to keep CAs accountable and that supports efficient, authentic, and fresh proofs of certificate status. To achieve this, RITM extends the concept of *authenticated dictionary* [29]. This structure allows a non-trusted *prover* to prove to a *verifier* that the queried object is (or is not) an element of a given dictionary created by a trusted third party. In our scenario, the RA is a prover, the client acts as a verifier, and the CA is a trusted third party. Every CA maintains a dictionary of its own revocations, while every RA stores copies of all the dictionaries. These dictionaries are constantly updated (by the CAs through the dissemination

insert (executed by a CA that revokes a certificate with serial number s_x):
 n : size of the dictionary with s_x appended, m : parameter chosen by CA

- 1) Insert s_x, n into the tree and rebuild it.
- 2) For random value v compute $H^m(v)$.
- 3) Return **signed root**: $\{root, n, H^m(v), time()\}_{K_{CA}^-}$ (1)

update (executed by an RA on input $s_x, \{root, n, H^m(v), t\}_{K_{CA}^-}$):

- 1) Verify the signature and timestamp t .
- 2) Insert s_x, n into the tree and rebuild it.
- 3) Accept the applied changes only if the newly-built root equals $root$ and n is a number of leaves in the new tree.

refresh (executed by a CA at least every Δ if there is no new revocation):
 t : timestamp of the latest signed root

- 1) $p = \lfloor (time() - t) / \Delta \rfloor$.
- 2) If $p < m$ return **freshness statement**: $H^{m-p}(v)$. (2)
- 3) If $p \geq m$ create a new signed root (see Eq. (1)) and go to step 1.

prove (executed by an RA on input s_x):

- 1) Produce presence/absence proof for s_x .
- 2) Return **revocation status**: $proof, \{root, n, H^m(v), t\}_{K_{CA}^-}, H^{m-p}(v)$, (3)
 where $H^{m-p}(v)$ is the latest (current) freshness statement.

Fig. 2: Interactions with an authenticated dictionary. Operations *insert* and *update* can be performed in batch (to add multiple revocations to the tree simultaneously).

network) with new revocations and metadata, which ensures that a given dictionary is fresh and consistent. Dictionaries are append-only (§VIII discusses this property and optimizations), and the proof generated by a prover: 1) can efficiently express presence/absence of a given revocation in/from a given dictionary, 2) is authentic, i.e., the verifier is sure that the statement was produced by the CA, and 3) is timestamped, so that a verifier can reject a stale proof. Moreover, authenticated dictionaries allow to keep CAs accountable, as they help to maintain a consistent view across the system. For instance, entities with a short proof can check whether their dictionary views are consistent. Consequently, all inconsistencies can be eventually detected and proven.

RITM implements authenticated dictionaries with hash trees (see §II). Every leaf of a tree is a serial number of the revoked certificate concatenated with the number of that revocation. For each dictionary, revocations are numbered consecutively, starting from 1. The numbering ensures that revocations are inserted into the tree in the correct order, which helps to maintain a consistent copy of the dictionary and detect attacks (such as revocation reordering). Leaves are sorted in lexicographical order (using their serial numbers), which allows to prove that a given leaf is or is not an element of the tree. We define the message names and the dictionary's interface in Fig. 2.

Dissemination. RITM can employ an existing CDN infrastructure as a dissemination network, and such a deployment brings many benefits. CDNs are efficient and robust in large-scale deployments [10], and we show that they provide an excellent environment for dissemination of revocation messages. As our study shows, a CDN-based deployment is also cost-effective (see §VII-C).

Whenever a CA wishes to revoke some certificate(s), it updates its local dictionary through the *insert* operation

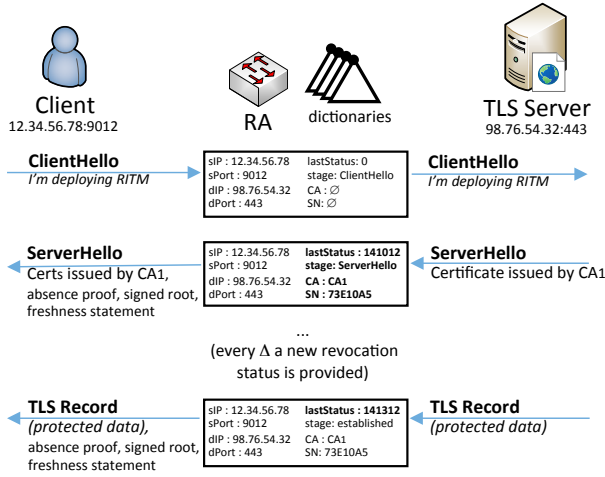


Fig. 3: An example of RITM-supported TLS connection.

(see Fig. 2). Then, the CA contacts the network's distribution point with the *revocation issuance* message, i.e., the revoked serial number(s) with a new signed root (e.g., at times t_0 and $t_0 + 3\Delta$ in Tab. I). The distribution point verifies this message and initiates the dissemination process by sending the message to edge servers. Eventually, all RAs obtain the message from the corresponding edge server and update their local copy of the CA's dictionary through the *update* operation (see §III).

When no new revocations are issued during a time period Δ , CAs are still obliged to keep their dictionaries fresh. In order to achieve this, every CA creates and disseminates a freshness statement (Eq. 2) periodically (e.g., at times $t_0 + \Delta$ and $t_0 + 2\Delta$ in Tab. I). Therefore, to keep dictionaries fresh when there is no new revocation within a time period Δ (which is a common case), CAs disseminate only a freshness statement that in practice is significantly shorter than a signed message. Moreover, due to hash chain properties (see §II), the message is authentic, as only a CA can create a valid freshness statement for a given time period. The dissemination process of a freshness statement is the same as before, and RAs, after verification, replace the freshness statement with the new one.

RITM provides a simple synchronization protocol to keep the dictionary copies of RAs correctly updated. As every revocation issuance message contains a value n that denotes the number of revocations that the CA has issued, an RA can easily detect whether his local copy is desynchronized from CA's original copy based on n . Whenever a desynchronization is detected, the RA contacts an edge server specifying the number of valid consecutive revocations it has observed.

Time	Revoked serial number	Disseminated message
$t = t_0$	s_a, s_b, s_c	$s_a, s_b, s_c, \{root, n, H^m(v), t\}_{K_{CA}^-}$
$t = t_0 + \Delta$	none	$H^{m-1}(v)$
$t = t_0 + 2\Delta$	none	$H^{m-2}(v)$
$t = t_0 + 3\Delta$	s_d	$s_d, \{root', n+1, H^m(v'), t\}_{K_{CA}^-}$

TABLE I: Example of messages disseminated over time.

The signed root with a timestamp (Eq. (1)) also guarantees integrity, as it is impossible to alter a dictionary without changing the root.

Validation. Our technique relies on the fact that the negotiation phase of TLS is performed in plaintext. Thus, RAs can detect new TLS connections and the corresponding server certificates that must be validated. A RITM-supported TLS connection is presented in Fig. 3, and we describe below the different steps in detail.

- 1) The TLS connection is initiated by a `ClientHello` message. The client sends it with a dedicated TLS extension [17], to inform potential RAs on the path that this TLS connection should be protected by RITM.
- 2) An RA on the path observes the traffic for TLS messages. Whenever it receives a `ClientHello` message with the RITM extension, it creates the following state:

$$\begin{aligned} sIP &= IP_c, \quad dIP = IP_s, \quad sPort = port_c, \quad dPort = port_s, \\ lastStatus &= 0, \quad stage = ClientHello, \quad CA = \emptyset, \quad SN = \emptyset, \end{aligned} \quad (4)$$

sIP , dIP , $sPort$, $dPort$ denote source/destination IP addresses and ports; $lastStatus$ is the latest time at which a revocation status was sent to the client; $stage$ describes the current stage of the TLS connection; CA allows the RA to identify the correct dictionary; and SN is the certificate's serial number. The RA identifies supported TLS connections using this state information.

When the state is created, the `ClientHello` message is passed to the next hop on the path towards the server.

- 3) The server responds with a `ServerHello` message followed by the server's certificate. (The server ignores the RITM's `ClientHello` extension.)
- 4) When an RA receives a `ServerHello` message that matches the previously-created state, it inspects the content of the message to obtain the server's certificate and to determine the CA that issued the certificate. Then, the RA executes the *prove* procedure (see Fig. 2) over the CA's dictionary. A revocation status (i.e., proof, signed root, and a freshness statement as in Eq. (3)) is appended to the `ServerHello` message and passed to the client. The RA updates the state from Eq. (4) by setting $lastStatus$ to the current time, $stage$ to `ServerHello`, CA to the CA identifier, and SN to the serial number of the server's certificate.
- 5) When the client receives a `ServerHello` message with a revocation status (Eq. (3)), he removes the status from the message (to not influence the TLS protocol), and a standard certificate-validation procedure is executed. Additionally, the client verifies the revocation status sent by the RA. The server's certificate is accepted when:

- a) it passes the standard validation,
- b) the revocation status contains a valid revocation absence proof (validated against a signed root), and
- c) the freshness statement is no older than 2Δ (see §V), i.e., for received revocation status as in Eq. (3), $H^{p'}(H^{m-p}(v))$ or $H^{p'+1}(H^{m-p}(v))$ equals $H^m(v)$, where $p' = \lfloor (time() - t) / \Delta \rfloor$.

- 6) Finally, when the TLS connection is accepted (i.e., when the server sends a `Finished` message), the RA updates the `stage` field to `established`. From this moment on, the client and the server communicate securely. Whenever the RA detects that Δ time has passed since it has received a revocation status for a supported connection (i.e., $time() - lastStatus \geq \Delta$), the RA uses the first TLS packet from the server to the client to piggyback a fresh revocation status (similarly as in step 4). After the fresh status is sent, the RA updates `lastStatus` to `time()`.
- 7) Whenever the client receives the packet with a revocation status, he checks the status, similarly as in step 5.

Steps 6 and 7 are conducted periodically, at least every Δ , and the connection is interrupted by the client, when a fresh absence proof is not provided. Whenever a supported connection is finished or timed out, the RA removes the corresponding state.

RITM supports two mechanisms of TLS resumption, namely *session identifiers* [17], and *session tickets* [37]. Although the TLS handshake for these modes is abbreviated, the presented mechanism is similar. We present RITM in a setting where the revocation status is provided to the client piggybacked on standard TLS traffic. However, in §VIII we discuss other implementation choices to achieve this functionality.

Consistency Checking. RITM enables RAs and clients to monitor CA actions related to certificate revocations. Through the construction of trusted dictionaries and validation logic, a CA has to provide freshness statements, which are short, unique, and unforgeable pieces of information about a dictionary's content. This message can be used directly by RAs to monitor the consistency of dictionaries. The goal is to ensure that a given party of the system has the same view of the dictionaries as the rest of the system. In order to achieve this property, an RA can periodically request a random edge server for its copy of the signed root. Only by comparing local and downloaded values, the RA can confirm that a copy stored by the edge server is identical to the RA's copy. Simply exchanging the latest signed root is enough to keep CAs accountable, as dictionaries are append-only, and as a violation from the append-only property is easily detectable.

Due to nature of the CDN environment an RA can detect and contact only a limited number of edge servers (usually, a list of the closest edge servers can be obtained via a DNS query). This limitation may decrease the effectiveness of the consistency checking procedure. To address this problem, a *map server* can be introduced. That server would store addresses of RAs (and optionally edge servers), so that they can communicate directly to exchange their current freshness statements. An alternative way is to deploy a gossip protocol (for example, as proposed by Chuat et al. [13]), where clients would exchange signed roots obtained during TLS connections, or to deploy dedicated parties to monitor CAs as proposed in log-based approaches [24].

Such procedures can protect from various possible attacks launched by a malicious CA (for details see §V). Moreover,

any detected misbehavior can be cryptographically proved and reported, for example to the software vendors.

IV. DEPLOYMENT MODELS

Incremental deployment is a challenging process for almost all new security technologies. The main challenges, associated with the deployment of RITM, are the following: 1) How can the RITM-supported clients be protected from downgrade attacks (in which the attacker convinces the client that there is no RA on a path towards the server)? 2) How can the new system be deployed in a backward-compatible manner? 3) How can RITM be efficiently deployed, i.e., so that a small number of RAs support all TLS traffic?

We present two deployment models for RITM, which we believe are feasible with current networks and that address the above-mentioned problems. The intuition behind the presented models is that RITM can work effectively when RAs are placed close to clients or servers, as it is easier to mitigate downgrade attacks and ensure that all TLS traffic passes through the RA. In §VIII we discuss other deployment strategies and aspects.

Close to the servers. Our first deployment model assumes that RAs are placed close to the servers, e.g., at the ingress point of a data center or a server farm. Typically, a load-balancer that distributes workload across multiple servers by inspecting the content of the packet is placed at the ingress point of a data center. For encrypted traffic (via TLS), the load-balancer additionally functions as a *TLS terminator*, which establishes and maintains TLS connections on behalf of the servers [16], to terminate TLS connections and inspect the decrypted traffic. RA functionality can be augmented to TLS terminators so that every TLS connection between a data center and its clients could benefit RITM. In this scenario, the augmented TLS terminators put an indication that the server side supports RITM within an extensions of TLS' `ServerHello`. This approach eliminates downgrade attacks on TLS connections since an adversary cannot undetectably alter the TLS handshake messages.

Close to the clients. Alternatively, RAs can be placed close to the clients, e.g., close to the gateways of clients' access networks. In fact, an RA can even be combined with the gateway, minimizing setup and deployment costs for the network operator that operates clients' access network. This approach provides two benefits: it protects client's connections to any TLS server regardless of deployment at the servers, and the network operator can place RAs in his network in such a way that all traffic in the network will traverse through a RA, hence, offering protection to the entire network. To this end, RAs can be either placed at the *choke points* of the network, or SDN-based solutions can be used [33].

However, the client-side deployment comes with a security vulnerability: it is susceptible to downgrade attack by an adversary that is within the client's access network and tunnels client's traffic to an external network so that the traffic bypasses a RA in the network. To avoid such downgrade attack, a network operator must provide clients with an authentic

information that his network supports RITM. This can be realized by secure network management protocols [18].

V. SECURITY ANALYSIS

Short Attack Window. Through a fast and robust dissemination and with an appropriate certificate-acceptance policy (i.e., clients may accept a certificate only with its fresh revocation status), RITM reduces the long attack windows associated to many revocation schemes. Effectively, the attack window is 2Δ . Although CAs publish and RAs download updates every Δ , they do not have to be synchronized. Due to the communication technique of CDNs (i.e., pull), a situation where an RA downloads an update just before the CA publishes new changes may occur. Hence, a *tolerance* parameter is required, and Δ is set as a value of this parameter.

The value of the Δ parameter is a trade-off between security and efficiency. However, as an RA is an online device connected to the dissemination network, we may expect even low values (tens of seconds) as feasible for a production setting (see §VII-B). Revocation checking in RITM is independent from server configuration, hence it is impossible to introduce a long attack window through a server-configuration parameter.

Race Condition. For an established TLS connection, client and RA perform periodic revocation checks. Such a design protects from a race condition in which a long-lived TLS connection with a server is established just before the corresponding certificate is revoked. In such a case, the client would unintentionally communicate with a non-trusted party until the connection is terminated. To the best of our knowledge, RITM is the first revocation system that protects from such an attack.

MITM and Blocking Attack. A MITM adversary, without the ability to bypass an RA cannot get any advantage. As it is assumed that the adversary cannot break the TLS protocol, he cannot undetectably modify the `ClientHello` message which carries RITM's extension. This message informs the RA on the path that the connection should be supported by RITM. Such an adversary can drop or delay status messages carried on TLS traffic but, according to RITM's validation policy (see §III), this would lead to a connection interruption.

Downgrade Attack. An adversary able to tunnel TLS traffic, in order to avoid an on-path RA, can try to launch a downgrading attack (as outlined in §IV), convincing a client that there is no RA on the path towards the server. This attack is generic and powerful against the majority of security protocols that are partially deployed.

To eliminate such an attack, the client must be provided with authentic information that the given connection is supported by an RA. Protection from this attack depends on the deployment scenario and we provide two feasible solutions. In the first deployment scenario RITM is deployed by TLS terminators (see §IV). In this case, the client would receive a deployment confirmation from a TLS terminator, within a `ServerHello` message. This message is itself protected by the TLS protocol, thus the adversary cannot forge it. As a consequence, the adversary is not able to launch a successful downgrade attack against RITM. Second, RITM is deployed in a network, which

can provide clients authentic information about RITM support. For instance, this information can be delivered through a secure equivalent of the DHCP protocol. After the bootstrapping, the client would expect a RITM protection for every TLS connection, so that even an adversary able to tunnel the traffic cannot bypass the scheme.

We do not discuss downgrade attacks against the TLS protocol itself, such as HTTPS stripping attacks [28], as this goes beyond the scope of this paper. However, we stress that this problem has been investigated in the past and some effective countermeasures are currently deployed [23], [20].

RA and Dissemination Network Compromise. Another advantage of RITM is that it requires a small trust base (i.e., only CAs are trusted). RITM is thus resilient to RA compromises. As the construction of authenticated dictionaries allows RA and CDN to be untrusted, an adversary that compromises an RA and CDN cannot get any benefit. Trusted dictionaries are created by CAs and, to forge a revocation status (see Fig. 2), the adversary would need to forge a digital signature, or break the hash function. Both conditions lead to the contradiction, as it is assumed that cryptographic primitives are secure. Moreover, an adversary with control over RAs cannot suppress any revocation message, as clients expect to see fresh revocation status messages periodically.

Misbehaving CA. Furthermore, RITM makes detectable a whole range of attacks that may be conducted by a misbehaving CA. The dictionaries are append-only, authenticated structures, and a signed root (Eq. 1) uniquely corresponds to a particular version of the dictionary. Therefore, it is possible to detect that a CA is misbehaving, for instance if the CA shows one version of the dictionary to part of the system, and another version to the rest. To prove that, it is enough to find two different signed roots with the same dictionary size. This can be realized in many ways (see §III). For example, RAs can randomly contact CDN edge servers or other RAs and compare their locally-stored statements with the newly-downloaded ones. It is sufficient to compare only the latest signed roots (with the same dictionary size), as dictionaries are append-only; therefore, whenever a misbehaving CA creates two versions of a dictionary, then the CA must constantly maintain these two different dictionaries (to satisfy update procedure from §III). Such a procedure guarantees that the system has a consistent view of all issued revocations. The append-only property and consecutive revocation numbers also protect from attacks such as revocation reordering, or revocation deletion. This substantially limits the possibilities of a misbehaving CA.

RITM can be also extended to the scenario where a compromised CA wishes to revoke its own certificate without introducing a collateral damage (i.e., without revoking certificates issued before the compromise). To handle such a case, a method presented by Szalachowski et al. [43] can be adopted.

More powerful adversaries. The attacks in which an adversary can capture a CA and another element of the system are especially dangerous. For instance, when an RA and a CA are compromised at the same time, then the adversary

can create a fake version of the CA’s dictionary (e.g., without a given revocation), and provide a proof to any client, using this fake dictionary. Although this scenario is unlikely, RITM can be extended to make even such an attack visible. As clients receive proofs accompanied with signed roots, they could contact edge servers, other RAs, or even other RITM-supported clients, to ensure that their view of the dictionary is consistent. The exact details of such a detection mechanism go beyond the scope of our current work, but our system can be for instance enhanced by gossip protocols [13].

Privacy. RITM preserves user privacy. Clients do not need to establish any extra connection for checking the status of a certificate and only the entities on the path towards the server know which clients contact which servers. In particular, CAs and the CDN network are not able to obtain this information.

VI. IMPLEMENTATION

To prove feasibility of RITM’s deployment we implemented it and the implementation includes 1) authenticated dictionaries, 2) an RA communication module (to contact the dissemination network and update dictionaries), 3) a deep packet inspection (DPI) module for RAs to analyze packets, identify new TLS connections, and keep track of the connection state, and 4) an RITM-supported client. RITM is implemented in Python (2.7.6), and we used Scapy with the Scapy-SSL/TLS module to realize DPI of TLS connections. We used the SHA-256 hash function, but we truncated its output to the first 20 bytes. Also, to optimize the bandwidth and computational overhead, we used the Ed25519 [12] signature scheme, which allows to obtain a digital signature of 64 bytes.

RAs are implemented as userspace daemons. All network traffic goes through an RA that verifies whether a packet belongs to the TLS handshake protocol or a previously established (and active) RITM-enabled connection. Then, the RA follows the different steps described in §III. For every TLS connection (after receiving a `ServerHello` message) RITM-supported TLS client expects a revocation status from the corresponding CA (delivered by an RA), which states that the certificate is not revoked. After this check passes, the client continues the standard TLS validation. For the communication between an RA and the CDN, we built a simple HTTP(S)-based API. Every Δ , each RA contacts an edge server via an HTTP GET request to pull new revocations and freshness statements. Our implementation of RAs can be adapted to operate with any CDN infrastructure.

VII. EVALUATION

A. Dataset used

We use a dataset collected by the *Internet Storm Center* [9], which contains the most complete list of CRLs (to the best of our knowledge), namely, 254 separate revocation lists. From these lists, we identify 1,381,992 unique revocations and an average of 5,440 revocations per CRL. Moreover, we identify that 3 bytes are used most frequently (32% of all revocations) as the size of serial numbers. Thus, we use 3-byte serial numbers throughout this analysis.




Fig. 4: Number of revocations issued between January 2014 and June 2015 with a focus on the Heartbleed peak.

To show how RITM reacts to extremely high revocation rates, we investigate revocations issued in April 2014. This time period is especially interesting because it includes the disclosure of the infamous *Heartbleed* [3] OpenSSL vulnerability. This bug allowed remote attackers to read protected server memory, including private keys associated with certificates. Consequently, thousands of certificates were revoked over a few days. The detailed analysis of this event is presented by Durumeric et al. [19] and Zhang et al. [46]. Fig. 4 (top graph) represents the number of revocations observed from January 2014. The Heartbleed event is visible as a sudden peak in the middle of April 2014. The highest revocation rate was observed on 16–17 April 2014 (bottom graph).

B. Speed of Dissemination

RITM strives to make revocations as readily accessible as possible, and one key criteria towards this goal is to minimize the time needed by RAs to download revocation messages. Using Amazon CloudFront [1] as an example CDN, we measure the time required for an RA to download revocation messages. Amongst many different CDN networks, we chose Amazon CloudFront because it allows to turn off content caching at the edge servers (by setting the caching interval to zero, i.e., TTL=0). This feature allows us to measure the worst-case scenario performance (in terms of latency), since the content needs to be fetched from the origin server for every request. Furthermore, since RAs in different geographical areas may experience different download latency, we take measurements from various locations. To this end, we use the Planetlab [14].

For measurements, we create revocation messages of different sizes, and upload them to Amazon CloudFront. More specifically, we test for five different revocation messages: only a freshness statement (without revocations), and revocation messages with 15K, 30K, 45K, and 60K revoked certificates. Then, from 80 Planetlab nodes from different geographical areas, we download the revocation messages and measure the time it took to download the messages. We repeat this experiment 10 times from each of the Planetlab nodes and for




Fig. 5: CDF of download times for five different revocation messages.

each of the five revocation messages. The measured download times are shown as CDFs in Fig. 5.

Our measurements show that even for a large number of revocation messages (e.g., the purple line with diamond markers) and even for the worst-case scenario (i.e., revocations are not cached), 90% of the Planetlab nodes took less than one second to download the revocations.

C. Cost of Dissemination

CDN operators charge content providers based on the amount of traffic that travels in and out of their network. In our case, a CA is a content provider that uploads the revocations and freshness statements and pays the CDN network for the traffic that RAs generate to download the data. In this section, using Amazon CloudFront as an example of CDN, we evaluate the price that a CA would need to pay to the CDN operator.

For this evaluation, we need a time series of the number of certificates revoked by a CA and the number and distribution of RAs deployed around the world. The distribution of RAs is necessary because Amazon CloudFront has different pricing rates for different locations [1].

For the former information, we use the largest CRL file that we could find (in terms of number of distinct revocations) [2]. It contains 339,557 revocation entries (almost 25% of all the revocations we gathered) with a total size of 7.5 MB. This case is extreme, as we determine the average number of revocations to be 5,440 per CRL.

There is no easy way to determine the number of RAs and their geographical distribution. One possible approach is to estimate the number of RAs with the number of IP addresses assigned to different parts of the world. However, such an estimation may be inaccurate since NATs are widely deployed. Instead, we use city population statistics; i.e., we estimate that the number of RAs is proportional to the population size and consider geographical locality at the granularity of cities. For the city population data, we use the MaxMind [4] dataset, which reports the population size as well as the latitude and longitude coordinates of the cities. The dataset has information for 2.3 billion people from 47,980 cities. Then, we determine the location of the edge servers that would serve the RAs in each city. Since we do not have access to computers in each of the city, we use the closest Planetlab node from each city to determine an edge server that would server an RA.




Fig. 6: Monthly bills for a CA using a CDN for disseminating revocation lists (where there are 10 clients per RA).

Clients per RA	$\Delta=10\text{ sec}$	$\Delta=1\text{ min}$	$\Delta=1\text{ h}$	$\Delta=1\text{ day}$
30	18.574	3.450	0.647	0.108
250	2.229	0.414	0.078	0.013
1,000	0.557	0.103	0.019	0.003

TABLE II: Average cost (in thousands of USD) in function of Δ and the number of clients per RA.

Our cost estimates are conservative: 1) we use the standard Amazon pricing, but when a CA negotiate with Amazon, the pricing should be much lower, 2) we assume that every RA serves only ten clients (thus there are 230 million RAs in total), 3) we analyze small values of Δ , and 4) the revocation list considered is the largest found. In reality, the monthly operating cost for the CA should be lower than our simulation, making CDN-based deployment even more attractive.

Fig. 6 shows the monthly bill that the CA would have to pay from 1 January 2014 to 1 August 2015, if they were to use RITM to distribute revocations. The simulation period includes the Heartbleed incident that occurred in April 2014 to account for any catastrophic event that results in a large number of revocations, and we show the results for four different values of Δ . We also estimate the average cost for the same CA, depending on how many clients a single RA handles and the results are presented in Tab. II. The values are given for small network (30 clients), large network (250 clients), and a middle-size corporate network (1,000 clients).

D. Overheads

Storage. Distribution points, RAs, and CAs are the only parties that store revocations in RITM. Neither clients nor servers need to store any revocation messages. This property makes RITM an ideal solution for hardware-limited clients (e.g., in the Internet of Things) that are not able to store all complete and synchronized revocation lists. Moreover, through RITM's design, the number of RAs should be substantially smaller than the number of TLS clients and servers.

With the above-mentioned dataset and setting, we estimate the storage overhead required by RAs to store the revocations. Conservatively, we assumed that all revocations (even for expired certificates) are held by an RA. We also assumed that dictionaries are implemented as described in §III. In such a setting, the storage overhead is slightly above 4 MB and the

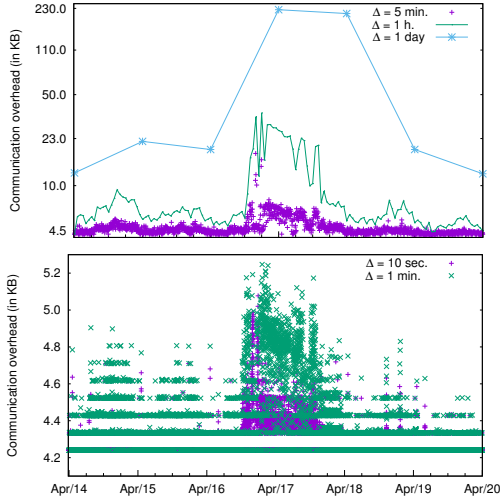


Fig. 7: Communication overhead depending on a revocation rate and the Δ parameter. Note that the top y-axis is in logarithmic scale.

memory required to build and keep all dictionaries is 36 MB (for instance, for 10 million revocations this overhead is 30 MB and 260 MB accordingly).

Communication. RITM introduces a negligible communication overhead to the TLS connection. The only additional information sent during the TLS handshake is a proof of a revocation status. The size of this proof is logarithmically proportional to the number of revocations in a given dictionary, but it also depends on the cryptographic primitives (in our case hashes are 20 bytes long). In our setting, a revocation status (Eq. 3) for an entry corresponding to the largest CRL that we observed would be 500-900 bytes. This short message is sent only during the connection establishment and then every Δ .

Communication overhead within the dissemination network is determined by 1) the revocation rate (as RAs have to receive every revocation), 2) the number of CAs (more precisely, the number of dictionaries), and 3) the Δ parameter. To estimate the required bandwidth, we analyze revocations issued during the week in which the Heartbleed vulnerability was disclosed. This period includes standard and extremely high revocation rates. Conservatively, we assume the number of dictionaries to be 254 (the number of CRLs we found), and we show results for Δ equals 10 seconds, 1 minute, 5 minutes, 1 hour, and 1 day. The obtained results are presented in Fig. 7, that depicts how much data a single RA must download every Δ .

As depicted, for a standard revocation rate, the required bandwidth is about 4 KB/ Δ . This overhead is mainly determined by the size of the freshness statements and number of dictionaries (as every dictionary has to be refreshed every Δ). For the highest peak observed, the dissemination requires below 5 KB/ Δ (for small Δ s), around 25 KB for Δ equals 1 hour, and about 230 KB when the RA is updated every day.

Computation. Using our implementation we evaluate the computational overhead introduced by RITM. Each test was executed 500 times on a machine with an Intel i7-4790 (3.60GHz) CPU, 16 GB of RAM, and running a 64-bit Linux.

Entity	Operation	Max.	Min.	Avg.
RA	TLS detection (DPI)	14.07	1.90	2.93
RA	Certificates parsing (DPI)	57.94	15.97	19.95
RA	Proof construction	156.16	31.00	67.17
Client	Proof validation	101.09	46.01	54.51
Client	Sig. and freshness valid.	380.04	159.03	197.27

TABLE III: Detailed processing time (in μ s).

First, we determine the time required to perform a dictionary update (which is periodically conducted by CAs and RAs). For instance, to insert (see Fig. 2) 1,000 new revocations, a CA needs to perform computations, which take 3.88, 2.75, and 2.93 ms, on maximum, minimum, and average, respectively. Similarly, for an RA to update a dictionary (see Fig. 2) with 1,000 new revocations and then verify the consistency of the dictionary and the authenticity of revocation issuance message, it takes 5.87, 2.62, and 2.84 ms.

We also analyze the latency introduced to the TLS connection. This latency is caused by computations executed by RAs and clients, and it is one of the most important metrics for connection establishment, as it has a direct influence on user reactions. The only additional latency in RITM is introduced by DPI (RA), proof construction (RA), and proof validation (client). We determine the computational time required for these operations. The detailed results are presented in Tab. III.

The RA needs to check whether each packet belongs to a TLS communication. With our implementation, this verification takes on average 2.93 μ s. RA must parse every `ServerHello` message to obtain the server's certificate. The computation time depends on the number of certificates in the chain (see §VIII), their length, used extensions, etc... For three certificates in the chain (the most common number [35]), such parsing takes on average 19.95 μ s. With the obtained certificate, the RA checks its revocation status, and returns the corresponding proof. This operation takes 67.17 μ s on average. Then, the revocation status obtained by the client must be verified. First, the client verifies whether the proof matches the obtained certificate and root (from the revocation status message). Then, the client verifies the signature of the signed root and corresponding freshness statement. These operations take on average 54.51 and 197.27 μ s, respectively. The overall overhead introduced on the client side is about 250 μ s. In a contrast to these results, a study showed that even an optimized TLS handshake takes about 30 ms (with network latency) [42]. Therefore, the overhead introduced by RITM to TLS connection establishment is less than 1%.

As the result, an RA can process more than 340,000 non-TLS packets per second and more than 50,000 RITM-supported TLS handshakes per second, on average. Clients can validate almost 4,000 revocation statuses per second. Thus, the overhead introduced by RITM should not be noticed by clients.

E. Comparison

We compare RITM with competing schemes (see §II) and present the results in Tab. IV. We consider the storage over-

head, the number of connections required by the scheme, and the achieved properties. For this comparison, we assume that the given revocation scheme is fully deployed. Then, we show the storage and the number of connections required to achieve a state in which an arbitrary client is able to establish a secure connection with an arbitrary server (i.e., the client must learn about the revocation status of the server's certificate). Log-based approaches are presented in their two deployment models (see §II). Global values show how the complete revocation list is replicated, and how many connections are required in total. Per-client values show how many entries a client must store, and how many messages it must obtain. The comparison demonstrates the scalability of RITM. Clients and servers neither store nor request for any messages. OCSP Stapling and log-based approaches (with server-driven deployment) have good performance results, but their deployment depends on the server's configuration, and they lack other desired properties.

F. Deployability

RITM has many advantages from a deployability point of view. An RA must provide a set of simple functionalities and can be cheaply implemented as a stand-alone device or as a part of another network device. Through the different deployment models that we present in §IV, we show that RITM's deployment can combine both, efficiency and security, and that RITM can be incrementally deployed. One RA can protect all TLS connections from/to a whole network. Therefore, a mass deployment is possible with a low operational effort. The deployment can be driven by a set of individual parties such as LAN operators, ISPs, or server farms operators. RITM is also backward-compatible as RAs are completely non-invasive for non-supported clients and protocols other than TLS.

VIII. DISCUSSION AND PRACTICAL CONSIDERATIONS

Bootstrapping CAs into RITM. Current CAs that wish to start RITM deployment could publish a signed manifest at pre-defined locations (e.g., /RITM.json). Such a manifest would contain the CDN address of a dictionary. Then, every RA would periodically check (e.g., once per week) whether a

CA started RITM deployment. Information about such a CA should be provided to the clients as well, and this could be realized through similar periodic checks or through software update (as the manifest is short).

Local Δ parameter. To simplify the presentation, we assumed that the Δ parameter was globally agreed on. However, to relax this assumption, each CA could express its own Δ parameter in a dedicated field of its certificate (or in a manifest as proposed above). Then, clients and RAs (who need CA certificates to validate server certificates or/and revocation statuses anyway) would know the correct Δ value, and operate accordingly.

Ever-growing dictionaries. We also assumed that each CA has a single append-only dictionary. Consequently, revocation entries (even for expired certificates) cannot be removed. However, this assumption can be relaxed as well. For instance, a CA can split all revocations into a few dictionaries. A similar practice is popular in maintaining CRLs (to reduce bandwidth overhead), and although it is not necessary for RITM, this can limit storage used on RAs. More precisely, a CA could maintain a few dictionaries at the same time, and every dictionary would be dedicated to all the certificates that expire before a given time. As “certificates issued after 1 April 2015 must have a validity period no greater than 39 months” (according to the CA/B Forum [6]), RAs could then regularly delete the dictionaries that correspond to expired certificates.

Certificate chains. We presented RITM in a setting where the validity proof is provided only for a server certificate. In practice, a chain of certificates is used. Although many revocation schemes ignore the revocation checking of CA certificates, RITM can be extended to support this feature. Instead of returning a single absence proof, an RA would return a proof for each certificate of the chain. The introduced overhead should not increase significantly, as certificate chains are usually short [35], and as the proof construction is an efficient operation (see §VII-D).

Multiple RAs. As RAs can be installed independently, it may happen that a single RITM-supported connection traverses multiple RAs. To maintain a constant overhead in such a situation, we require that an RA adds a revocation status only when it is missing, and replaces a revocation status only if its own version of the dictionary is more recent. Whenever an RA sees a revocation message sent by another RA, it can verify whether their view of the dictionary is consistent, by simply comparing the two versions of the signed roots.

RA-to-client communication. One challenge in implementing RITM is to reliably transfer a revocation status from an RA to a client along with a `ServerHello` message. Fortunately, middleboxes have been studied for years, and we can benefit from previous works and well-known techniques. We distinguish three possible methods:

- 1) The status is piggybacked on the native TLS message. Since, the payload of the TCP packet (i.e., the packet that carries the altered TLS message) must be extended, the RA must adjust the sequence numbers of the TCP session. Then, the RA must also indicate (e.g., through a dedicated TLS `Content Type`) that client should handle

Method	Storage (global)	Storage (client)	Conn. (global)	Conn. (client)	Violated properties
CRL	$n_{rev} \times (n_{cl} + 1)$	n_{rev}	$n_{cl} \times n_{ca}$	n_{ca}	I, P, E, T
CRLSet ^a	$n_{rev} \times (n_{cl} + 1)$	n_{rev}	n_{cl}	1	I, E, T
OCSP	n_{rev}	0	$n_{cl} \times n_s$	n_s	I, P, E, T
OCSP ^b	$n_{rev} + n_s$	0	n_s	0	I, S, T
Log ^c	n_{rev}	0	$n_{cl} \times n_s$	n_s	I, P, E
Log ^d	n_{rev}	0	n_s	0	I, S
RevCast ^e	$n_{rev} \times (n_{cl} + 1)$	n_{rev}	n_{cl}	n_{rev}	E, T
RITM	$n_{rev} \times (n_{ra} + 1)$	0	n_{ca}	0	-

TABLE IV: Comparison of revocation mechanisms in terms of properties and overhead. $n_s, n_{ca}, n_{ra}, n_{cl}, n_{rev}$: number of servers, CAs, RAs, clients, and revocations, respectively. $n_{ca} \ll n_{ra} < n_s \ll n_{cl}$.

I: near-instant revocation P: privacy E: efficiency and scalability T: transparency and accountability S: server changes not required.

^aCRLSets contain a limited number of revocations. ^bOCSP Stapling.

^cClient-driven approaches. ^dServer-driven approaches. ^eRevCast uses radio broadcast for dissemination.

the TLS message differently. Otherwise TLS would refuse the connection, as the TLS handshake was modified.

- 2) The client opens a dedicated port, to which RA sends RITM messages. Though, such a design is easy to realize, it does not allow to contact clients placed behind a NAT.
- 3) Instead of using a dedicated port, the port that the client is using for the TLS communication is used to send the revocation status. This approach solves the NAT problem, and does not influence the TLS protocol directly, but RAs still have to adjust the TCP state for supported sessions.

Future work. We will investigate the possible interactions between RITM and novel PKI enhancements [44], [43]. We will also analyze some aspects of the incremental deployment of technologies such as RITM. Other topics of interest include a study on ensuring the consistency between RAs and clients, and applying novel deployment models [40], where RITM can be implemented as a service.

IX. CONCLUSION

In this paper, we present RITM, a framework that aims to address the revocation problem through distinctive properties of middleboxes and content delivery networks. RITM takes advantage of an existing infrastructure to disseminate revocation messages, and moves certificate revocation lists to the middleboxes. Such a combination results in a scheme that satisfies all the requirements we identified. In particular, we believe that RITM is the first scheme that can provide near-instantaneous revocation in a real-world deployment. Among other advantages, RITM prepares us for future demands originating from, e.g., pervasive encryption or the Internet of Things, in which countless hardware-constrained devices would need to communicate securely.

REFERENCES

- [1] Amazon CloudFront. <https://aws.amazon.com/cloudfront>.
- [2] CAcert's CRL. <http://crl.cacert.org/revoked.crl>.
- [3] The Heartbleed bug. <http://heartbleed.com>.
- [4] MaxMind. <https://www.maxmind.com>.
- [5] Mozilla's revocation plan. <https://wiki.mozilla.org/CA:RevocationPlan>.
- [6] Baseline requirements for the issuance and management of publicly-trusted certificates. https://cabforum.org/wp-content/uploads/Baseline_Requirements_V1_1_6.pdf, 2014. CA/Browser Forum.
- [7] An evaluation of the effectiveness of Chrome's CRLSets. <https://www.grc.com/revocation/crlsets.htm>, 2014.
- [8] Let's Encrypt. <https://letsencrypt.org/>, 2014.
- [9] SSL CRL activity. <https://isc.sans.edu/crls.html>, 2015. SANS Internet Storm Center.
- [10] M. Al-Fares, K. Elmeleegy, B. Reed, and I. Gashinsky. Overclocking the Yahoo! CDN for faster web page loads. In *ACM IMC*, 2011.
- [11] D. Basin, C. Cremers, T. H.-J. Kim, A. Perrig, R. Sasse, and P. Szalachowski. ARPKI: Attack resilient public-key infrastructure. In *ACM CCS*, 2014.
- [12] D. J. Bernstein, N. Duif, T. Lange, P. Schwabe, and B.-Y. Yang. High-speed high-security signatures. *Journal of Cryptographic Engineering*, 2012.
- [13] L. Chuat, P. Szalachowski, A. Perrig, B. Laurie, and E. Messeri. Efficient gossip protocols for verifying the consistency of certificate logs. In *IEEE CNS*, 2015.
- [14] B. Chun, D. Culler, T. Roscoe, A. Bavier, L. Peterson, M. Wawrzoniak, and M. Bowman. PlanetLab: an overlay testbed for broad-coverage services. *ACM SIGCOMM Computer Communication Review*, 2003.
- [15] D. Cooper, S. Santesson, S. Farrell, S. Boeyen, R. Housley, and W. Polk. Internet X.509 public key infrastructure certificate and certificate revocation list (CRL) profile. RFC 5280, May 2008.
- [16] A. Delignat-Lavaud and K. Bhargavan. Network-based origin confusion attacks against HTTPS virtual hosting. In *WWW*, 2015.
- [17] T. Dierks and E. Rescorla. The transport layer security (TLS) protocol version 1.2. RFC 5246, August 2008.
- [18] R. Droms and W. Arbaugh. Authentication for DHCP messages. RFC 3118, June 2001.
- [19] Z. Durumeric, J. Kasten, D. Adrian, J. A. Halderman, M. Bailey, F. Li, N. Weaver, J. Amann, J. Beekman, and M. Payer. The matter of Heartbleed. In *ACM IMC*, 2014.
- [20] C. Evans, C. Palmer, and R. Sleevi. Public key pinning extension for HTTP. RFC 7469, April 2015.
- [21] N. Gruschka, L. L. Iacono, and C. Sorge. Analysis of the current state in website certificate validation. *Security and Communication Networks*, 2014.
- [22] P. Gutmann. How to build a PKI that works. In *PKI R&D Workshop*, 2004.
- [23] J. Hodges, C. Jackson, and A. Barth. HTTP strict transport security (HSTS). RFC 6797, November 2012.
- [24] T. H.-J. Kim, L.-S. Huang, A. Perrig, C. Jackson, and V. Gligor. Accountable key infrastructure (AKI): A proposal for a public-key validation infrastructure. In *WWW*, 2013.
- [25] A. Langley. Revocation checking and Chrome's CRL. <https://www.imperialviolet.org/2012/02/05/crlsets.html>, 2012.
- [26] B. Laurie and E. Kasper. Revocation transparency, 2012.
- [27] Y. Liu, W. Tome, L. Zhang, D. Choffnes, D. Levin, B. Maggs, A. Mislove, A. Schulman, and C. Wilson. An end-to-end measurement of certificate revocation in the web's PKI. In *ACM IMC*, 2015.
- [28] M. Marlinspike. More tricks for defeating SSL in practice. *Black Hat USA*, 2009.
- [29] M. Naor and K. Nissim. Certificate revocation and certificate update. *IEEE Journal on Selected Areas in Communications*, 2000.
- [30] Y. Petersen. The transport layer security (TLS) multiple certificate status request extension. RFC 6961, June 2013.
- [31] M. Prince. The hidden costs of Heartbleed. <https://blog.cloudflare.com/the-hard-costs-of-heartbleed/>, 2014.
- [32] M. Prince. Universal SSL. <http://blog.cloudflare.com/introducing-universal-ssl/>, 2014.
- [33] Z. A. Qazi, C.-C. Tu, L. Chiang, R. Miao, V. Sekar, and M. Yu. SIMPLIFYING middlebox policy enforcement using SDN. In *ACM SIGCOMM*, 2013.
- [34] E. Rescorla. The transport layer security (tls) protocol version 1.3. 2016.
- [35] I. Ristic. State of SSL. *InfoSec World*, 2011.
- [36] R. L. Rivest. Can we eliminate certificate revocation lists? In *Financial Cryptography*, 1998.
- [37] J. Salowey, H. Zhou, P. Eronen, and H. Tschofenig. Transport layer security (TLS) session resumption without server-side state. RFC 5077, January 2008.
- [38] S. Santesson, M. Myers, R. Ankney, A. Malpani, S. Galperin, and C. Adams. X.509 Internet public key infrastructure online certificate status protocol – OCSP. RFC 6960, June 2013.
- [39] A. Schulman, D. Levin, and N. Spring. RevCast: Fast, private certificate revocation over FM radio. In *ACM CCS*, 2014.
- [40] J. Sherry, S. Hasan, C. Scott, A. Krishnamurthy, S. Ratnasamy, and V. Sekar. Making middleboxes someone else's problem: Network processing as a cloud service. In *ACM SIGCOMM*, 2012.
- [41] SSL Labs. SSL Pulse. <https://www.trustworthyinternet.org/ssl-pulse/>, 2012.
- [42] E. Stark, L.-S. Huang, D. Israni, C. Jackson, and D. Boneh. The case for prefetching and prevalidating TLS server certificates. In *NDSS*, 2012.
- [43] P. Szalachowski, L. Chuat, and A. Perrig. PKI safety net (PKISN): Addressing the too-big-to-be-revoked problem of the TLS ecosystem. In *IEEE EuroS&P*, 2016.
- [44] P. Szalachowski, S. Matsumoto, and A. Perrig. PoliCert: Secure and flexible TLS certificate management. In *ACM CCS*, 2014.
- [45] E. Topalovic, B. Saeta, L.-S. Huang, C. Jackson, and D. Boneh. Towards short-lived certificates. *Web 2.0 Security and Privacy*, 2012.
- [46] L. Zhang, D. Choffnes, D. Levin, T. Dumitras, A. Mislove, A. Schulman, and C. Wilson. Analysis of SSL certificate reissues and revocations in the wake of Heartbleed. In *ACM IMC*, 2014.